

•  
•  
•  
•  
•  
•  
•  
•  
•  
•  
•

# XML Transactions



•  
•  
•

# XML as Database

- Multiple Users
  - 1~1000, or more?
  - Querying and updating occur simultaneously
- Transaction Management
  - Atomicity of query and update operations
    - All-or-nothing execution
  - Consistency and Concurrency Control
    - Locking system

- 
- 
- 

# XQuery Update

- XQuery Update Facility 1.0
  - <http://www.w3.org/TR/2008/CR-xquery-update-10-20080801/>
- Primitive Updates: insert, delete, replace, rename
  - Extensions to other expr: FLWOR, TypeSwitch, ...

- 
- 
- 

# Examples

- delete nodes `//book[@year lt 1968]`
- insert node `<author/>` into `//book[@ISBN eq "34556"]`
- for `$x` in `//book`  
    where `$x/year lt 2000` and `$x/price gt 100`  
    return replace value of node `$x/price`  
    with `$x/price-0.3*$x/price`
- if (`$book/price gt 200`) then  
    rename node `$book` as "expensive-book"

Update statements work on "node" or "nodes"

- 
- 
- 

# Insert

- Insert a new element into a document
- **insert node** UpdateContent **into** TargetNode
  - UpdateContent: any sequence of items (nodes, values)
  - TargetNode: Exactly one document or element
    - Otherwise ERROR
- Optionally, specify if to insert at the beginning or end
  - as last: Content becomes first child of Target
  - as first: Content becomes last child of Target
  - No position: no fixed position (honor other first/last inserts)
- Nodes in Content assume a new Id.

- 
- 
- 

# Insert

- Insert new book in the library  
**insert node** <book> <title>Die wilde Wutz</title> </book>  
**into** document("www.uni-bib.de")//bib
- Insert new book at the beginning of the library  
**insert node** <book> <title>Die wilde Wutz</title> </book>  
**as first into** document("www.uni-bib.de")//bib
- Insert new attribute into an element  
**insert node** (attribute age { 13 })  
**into** document("persons.xml")//person[@name = "KD"]

- 
- 
- 

# Insert

- Insert at a particular point in the document
- **insert node** UpdateContent (**after** | **before**) TargetNode
  - UpdateContent: No attributes allowed!
  - TargetNode: One Element, Comment or PI.
    - Otherwise ERROR
    - Must have parent
- Specify whether before or behind target
  - Before vs. After
- Nodes in Content assume new identity

Add an author to a book

**insert node** <author author>Florescu Florescu</ </author author>  
**before** // article[title = "XL"]/ [author . = " Vincent"]

- 
- 
- 

# Delete

- Deletes nodes from a document
- **delete (node | nodes) TargetNodes**
  - TargetNodes: Any sequence of nodes
- Delete papers.  
**delete node** //article[header/keyword = "XML"]



- 
- 
- 

# Replace

- Replace a node  
**replace node** TargetNode **with** UpdateContent
- Replace the content of a node  
**replace value of node** TargetNode **with** UpdateContent
  - TargetNode: One node (with Id)
  - UpdateContent: Any sequence of items
  - Keeps the node ID of TargetNode

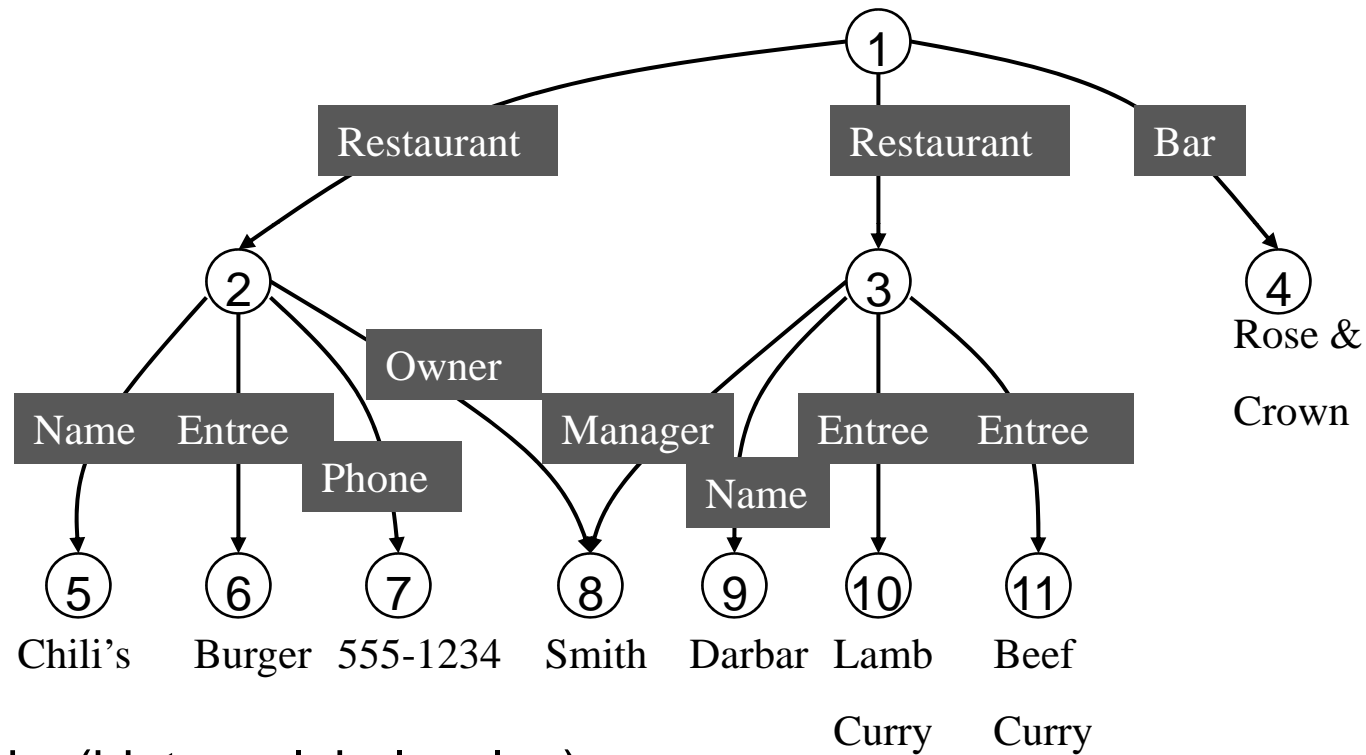
- 
- 
- 

# Concurrency Control

- Concurrency control is a method used to ensure that database transactions are executed in a safe manner.
- ACID properties
  - **A** tomicity: All actions in the transaction happen, or none happen.
  - **C** onsistency: If each transaction is consistent, and the DB starts consistent, it ends up consistent.
  - **I** solation: Execution of one transaction is isolated from that (the effect) of other transactions.
  - **D** urability: If a transaction commits, its effects persist.

•  
•  
•

# Dataguide model

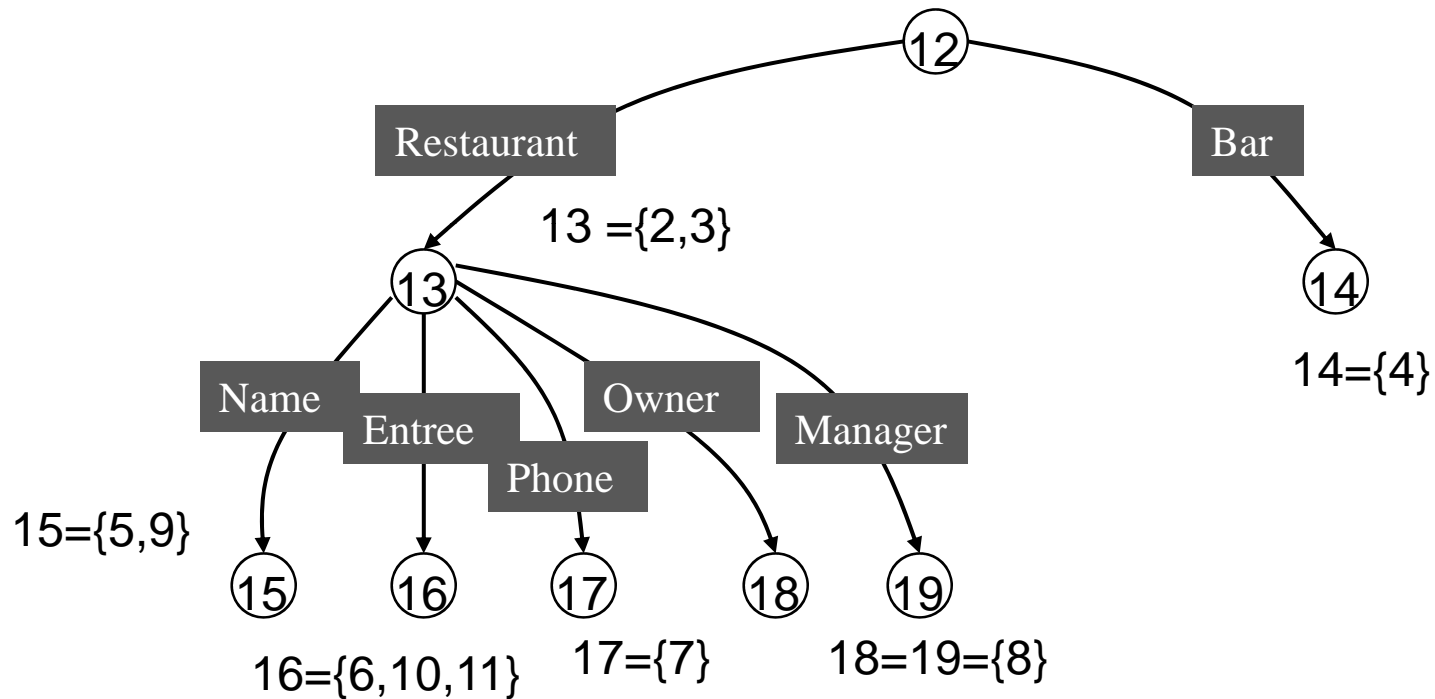


nodes(id, type, label, value)

child( child-id, parent-id)


•  
•  
•

# Dataguide model



•  
•  
•

# Data-value Conflict



Access Path  $L=/a/b/c$

T1: R(L)  $r(a)r(b)r(c)$

T2: W(L)  $r(a)r(b)w(c)$


$r(a)r(b) \ r(a)r(b) \ w(c) \ r(c) \ \text{abort}$

•  
•  
•

# Structural Conflict

L1=/a/d/e

L2=/a/d




T1: Append(L1,v) - r(a)r(d)r(e)w(v)w(e)

T2: Remove(L2) - r(a)r(d)w(d)

r(a) r(a) r(d)r(e)w(v) r(d)w(d) w(e)

•  
•  
•

## Element-order Conflict



$L = /a/g/f$

T1: Append(L, u)     $r(a)r(g)r(f)w(u)w(f)$

T2: Append(L, v)     $r(a)r(g)r(f)w(v)w(f)$

What is a possible problem?

•  
•  
•

## Data-value Conflict



L1=/a/c/d

L2=/a/b/d

T1: Read(L1)

T2: Write(L2)

index



Problem in the index  
tree, common D node



- 
- 
- 

# Structural Conflict and Index Node Split



L1=/a/b/d

L2=/a/b

T1: Append(L1, v)

T2: Remove(L2) Rename(F)

index



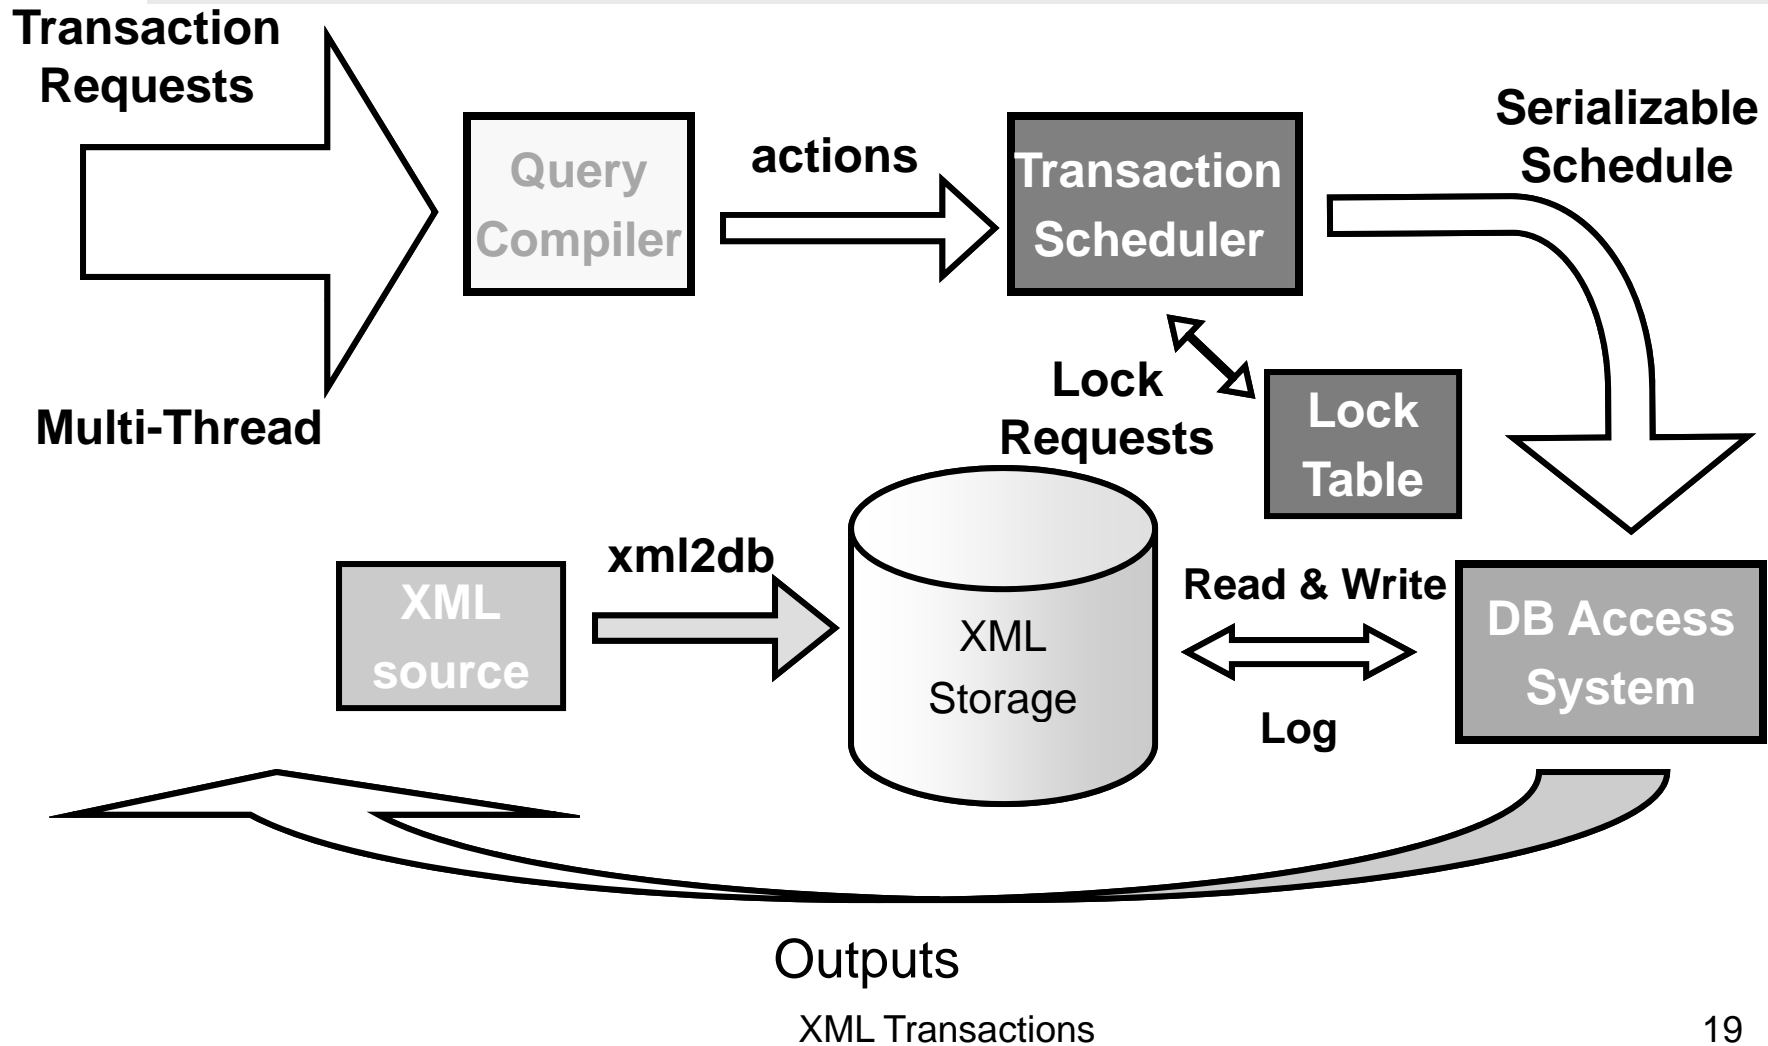
- 
- 
- 

# Xerial

- [www.xerial.org](http://www.xerial.org)
  - The name comes after XML and serializable
  - Taro L. Saito. On concurrency and updatability of XML databases. Master thesis submitted to Department of Computer Science, Graduate School of Information Science and Technology, University of Tokyo, January 2004
- Transactional Database for XML
  - Concurrent Transactions
    - Serializable schedule
  - Recoverability
    - Handling transaction aborts and system failures
  - Updating XML
    - Node insertion, deletion, modification, etc.
  - Transaction Language
    - Query and update notations

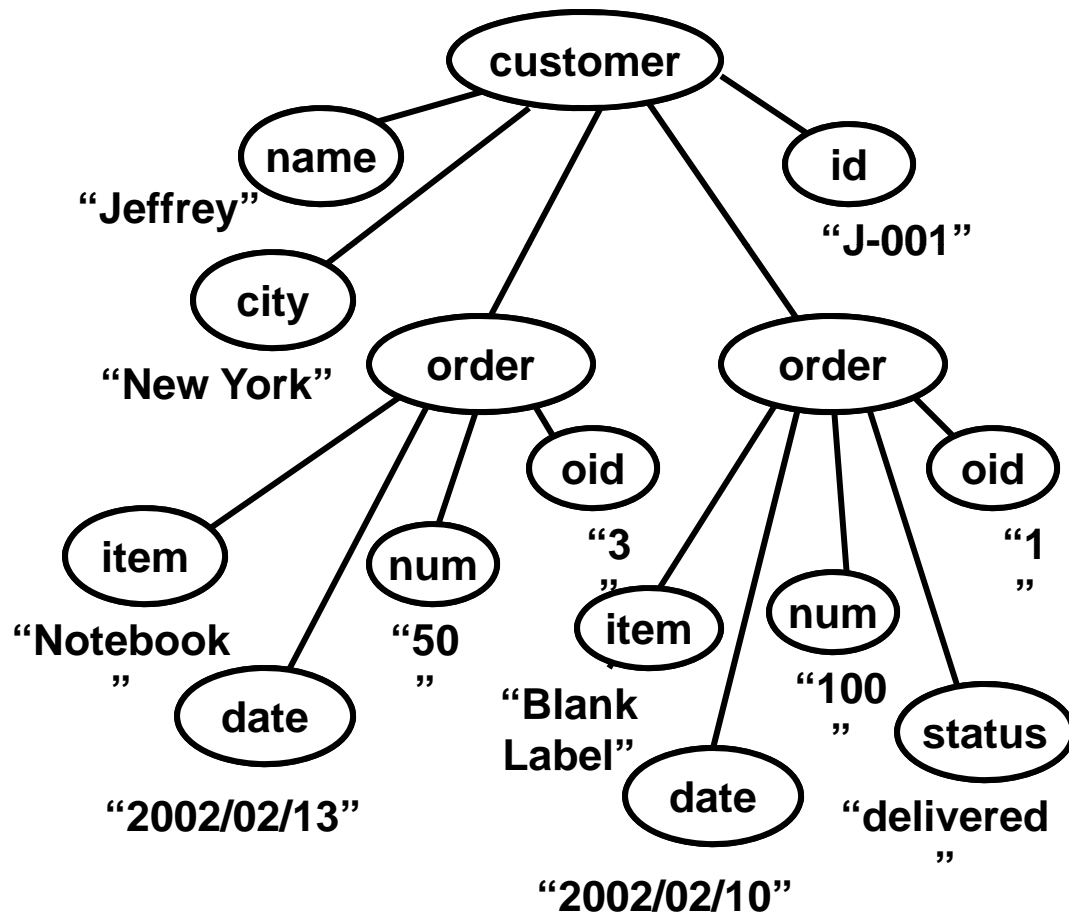
•  
•  
•

# Xerial Overview



•  
•  
•

# Data Model

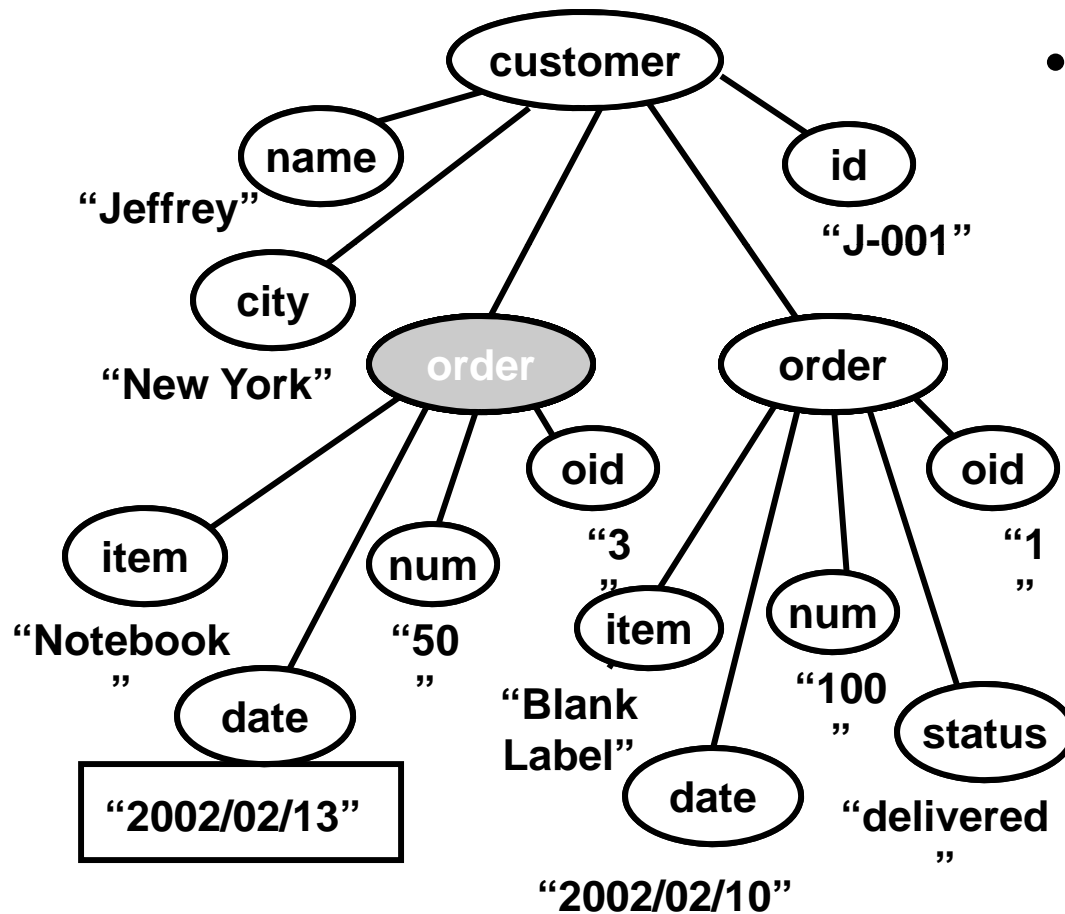


```

<customer id="J-001">
  <name> Jeffrey </name>
  <city> New York </city>
  <order oid="3">
    <item> Notebook </item>
    <date> 2002/02/11 </date>
    <num> 50 </num>
  </order>
  <order oid="1">
    <item> Blank Label </item>
    <date> 2002/02/10 </date>
    <num> 100 </num>
    <status> delivered </status>
  </order>
</customer>
  
```

•  
•  
•

# Querying XML



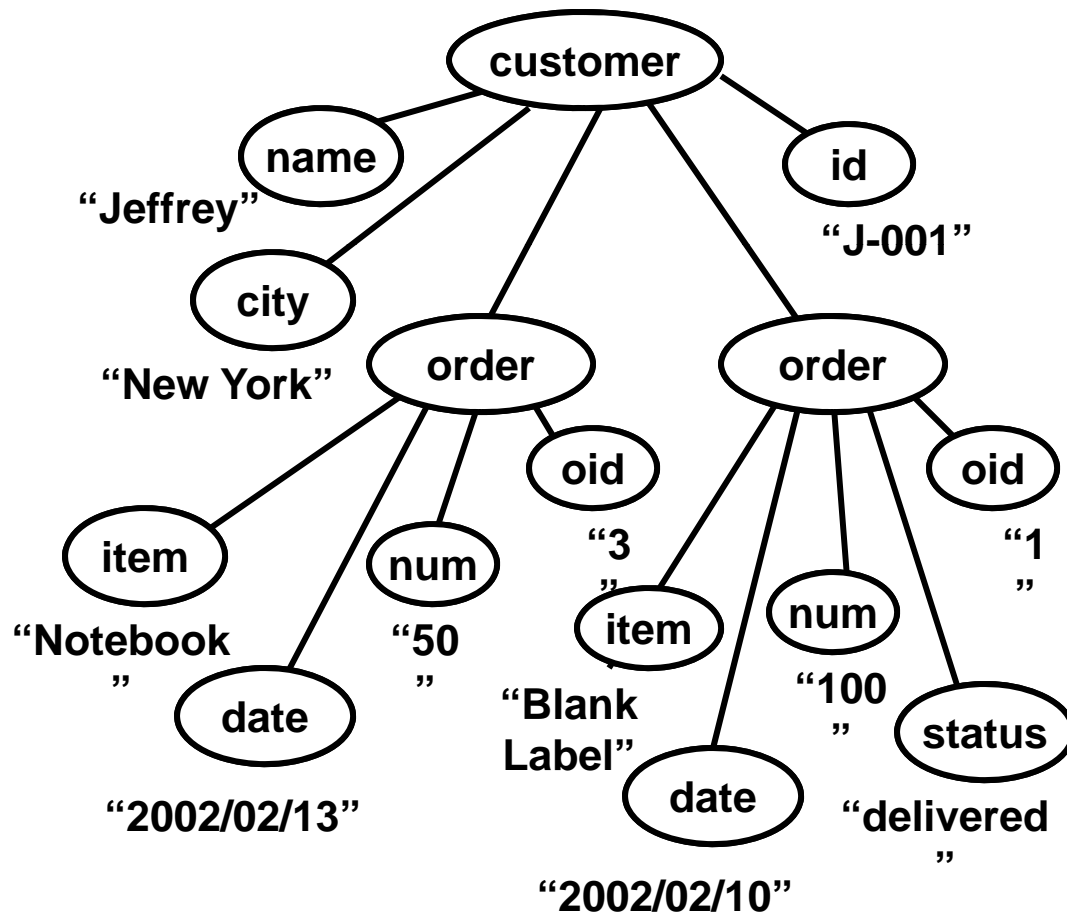
- **XQuery**

- W3C standard
- Query Language for XML
- Use of **Path expressions**
- Bind elements to a **variable**

```
FOR $x IN /customer/order
WHERE $x/date = "2002/02/13"
```

•  
•  
•

# Locks for Tree-Structure



- **Subtree Level Locking**
  - Query to entire subtree is frequent in XML
  - Reduce the number of locks
- **Performance Factor**
  - The number of locks
    - Load of lock manager
  - Granularity of locks
    - Concurrency

•  
•  
•

# Locks

	IS	IX	S	X
IS	Yes	Yes	Yes	No
IX	Yes	Yes	No	No
S	Yes	No	Yes	No
X	No	No	No	No

## Compatibility Matrix

### Ordinal Locks


- *S* Shared Lock (read)
- *X* Exclusive Lock (write)

### Warnings

- *IS* Intention to Share
- *IX* Intention to Exclusive

•  
•  
•

# Warning Protocol




- Jim Gray et al, 1975.
- Original Rules
  - All transactions must enter from the root
  - To place a lock or warning on any element, we must hold a *warning* on its parent
  - Never remove a lock or warning unless we hold no locks or warnings on its children



•  
•  
•

# Warning Protocol for XML



- **Extension**

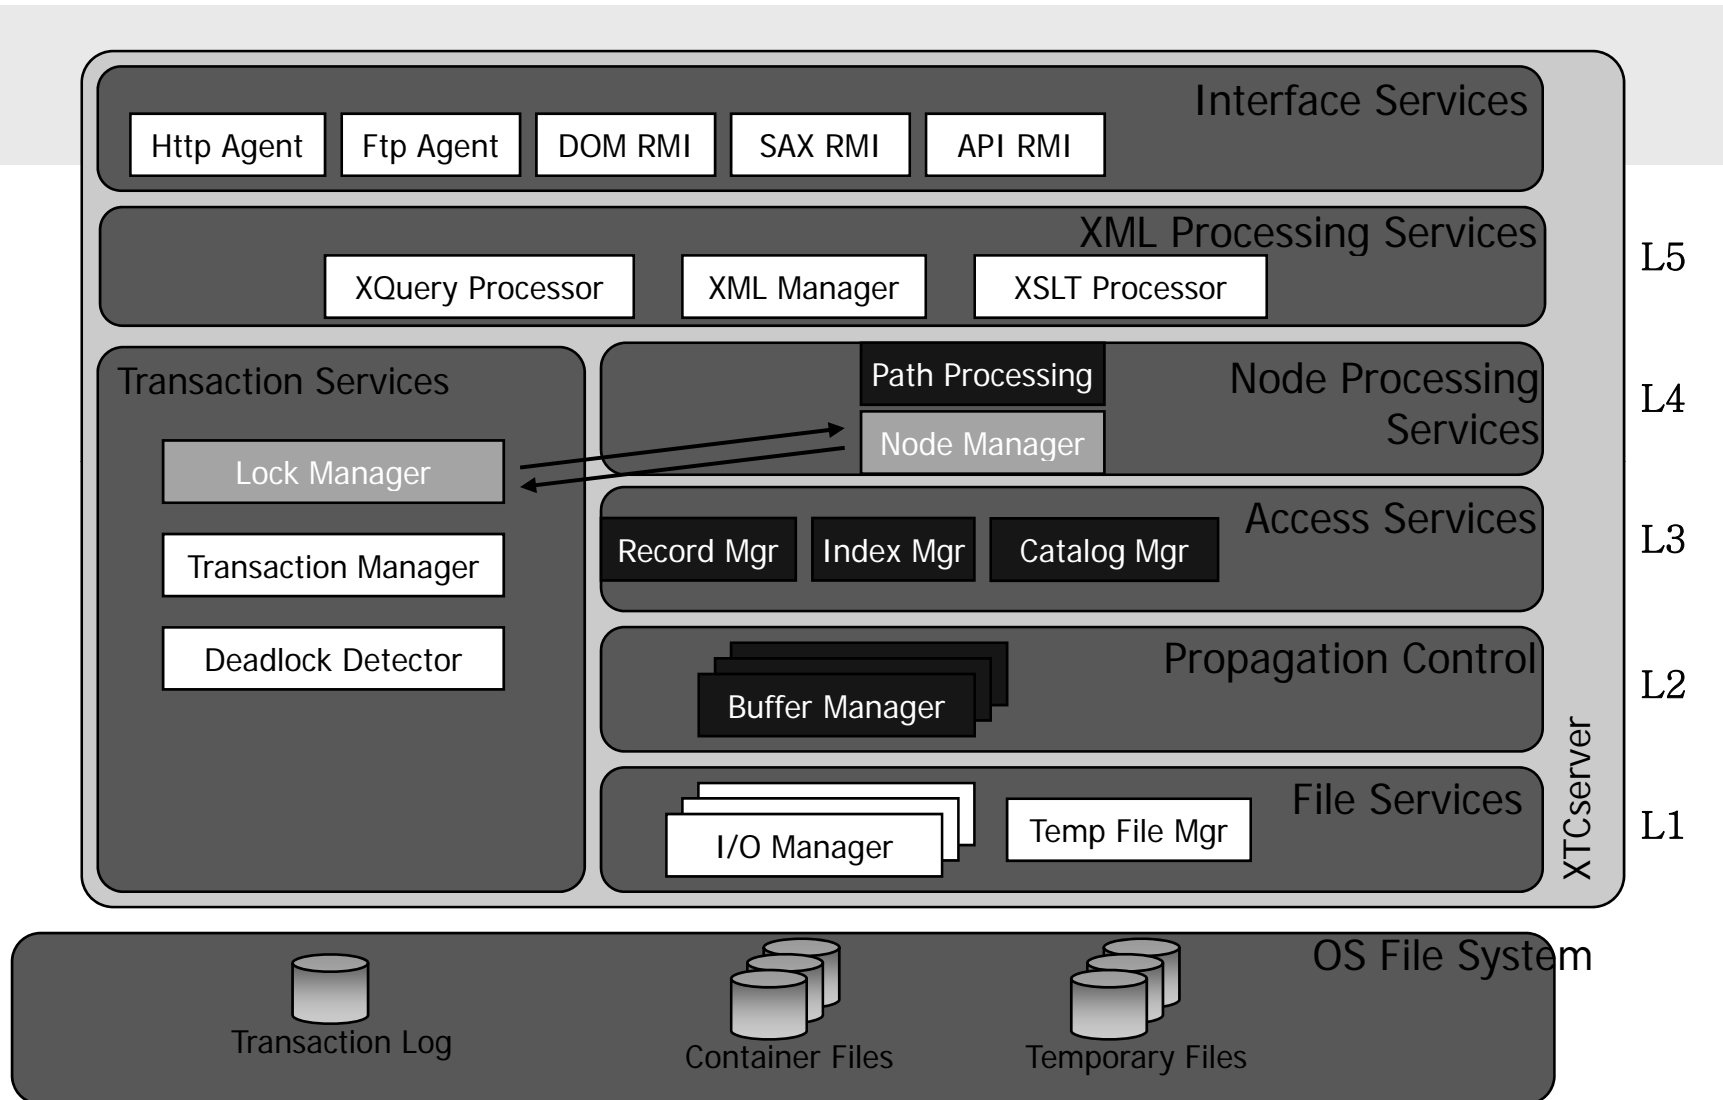
- When we insert or delete nodes, we must obtain *X* lock on the *parent* of the destination
- Until we place a warning on a node, we cannot trace its pointers to the children
- A transaction never releases locks or warnings until it finishes
  - **2 phase locking**

- 
- 
- 

# XTC

- XTC is used as a test vehicle for empirical DB research
  - effectiveness of XML concurrency control
    - fine-granular locking on nodes and edges
    - meta-synchronization allows comparison of different compatibilities
    - taDOM\* protocols
      - multiplicity of lock modes
      - intention locks are important
      - indexed document access is frequent
      - ancestor path locking without accessing the storage engine desirable
- References
  - Michael P. Haustein, Theo Haerder: Twig Query Processing Under Concurrent Updates, in ICDE. 2006.
  - Michael P. Haustein, Contest of XML Lock Protocols, VLDB 2006 ([www.vldb.org/conf/2006/p1069-haustein.pdf](http://www.vldb.org/conf/2006/p1069-haustein.pdf))

# XTC – Architectural Overview



determination of ancestor node IDs are of outmost importance  
for any locking protocol XML Transactions

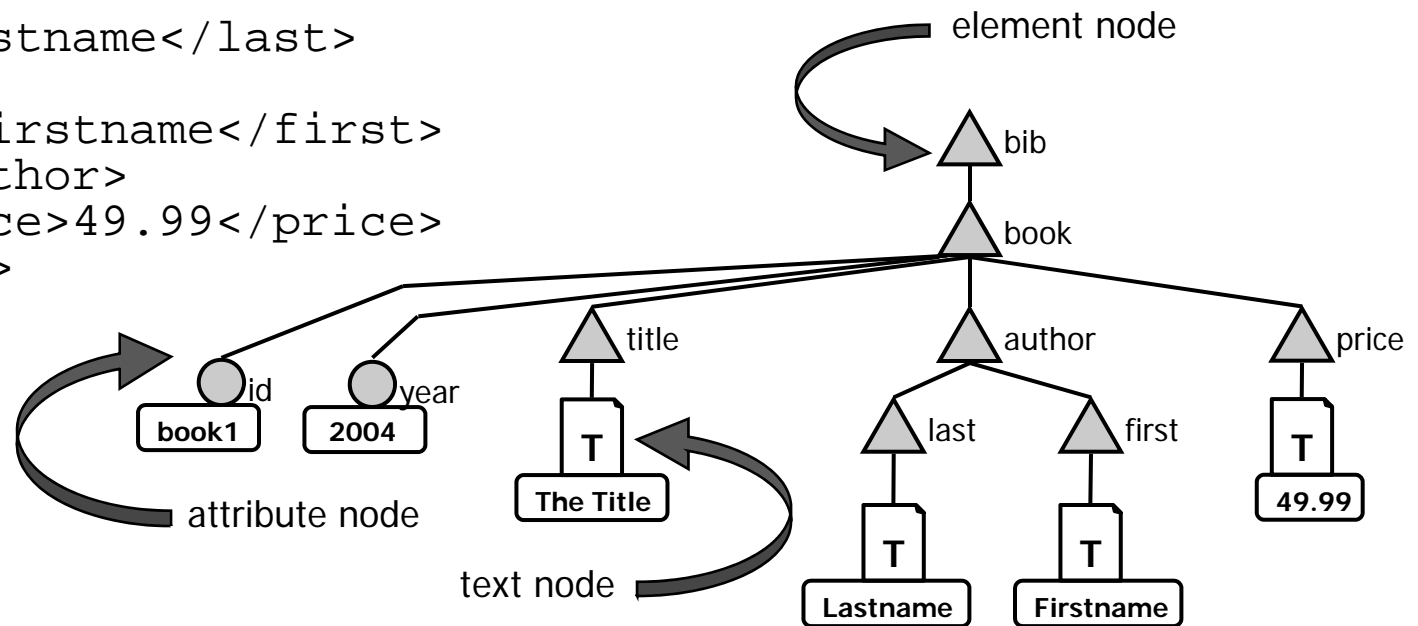
•  
•  
•

# XML Document

```
<?xml version="1.0"?>
  <bib>
    <book year="2004"
    id="book1">
      <title>The
Title</title>
      <author>

<last>Lastname</last>

<first>Firstname</first>
      </author>
      <price>49.99</price>
    </book>
  </bib>
```



- 
- 
- 

## Tailored Node Locks for XML – taDOM2

Read locks

lock	effect
NR (node read)	locks only a context node
LR (level read)	read lock on a context node and all direct-child nodes

Write locks

lock	effect
CX (child excl.)	indicates existence of a write lock on a direct child node
SU (update option)	read lock for intended update operations on an entire subtree

- 
- 
- 

## Tailored Node Locks for XML – taDOM2

Compatibility matrix

	-	IR	NR	LR	SR	IX	CX	SU	X
IR									
NR	+	+	+	+	+	+	+	-	-
LR	+	+	+	+	+	+	-	-	-
SR									
IX									
CX	+	+	+	-	-	+	+	-	-
SU	+	+	+	+	+	-	-	-	-
X									

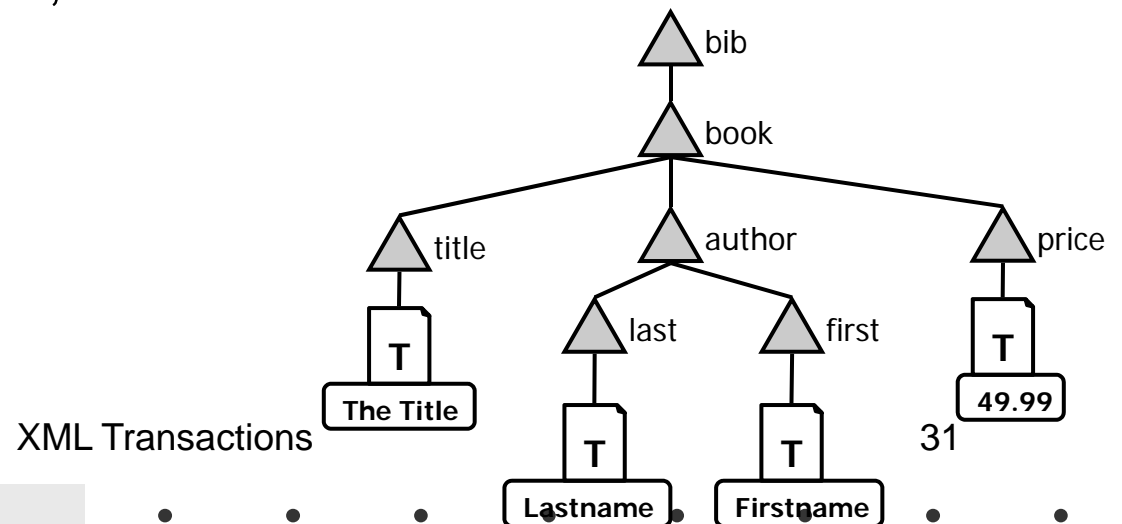
XML Transactions

30

- 
- 
- 

# Node Locks

- Node read lock (NR)
  - requires IR locks on the ancestor path
- Level read lock (LR)
  - requested for reading the context node and all nodes located at the level below (all direct-child nodes)
- Child exclusive lock (CX)
  - indicates an X lock on a child
  - defined, in addition to IX, to detect conflicts with LR



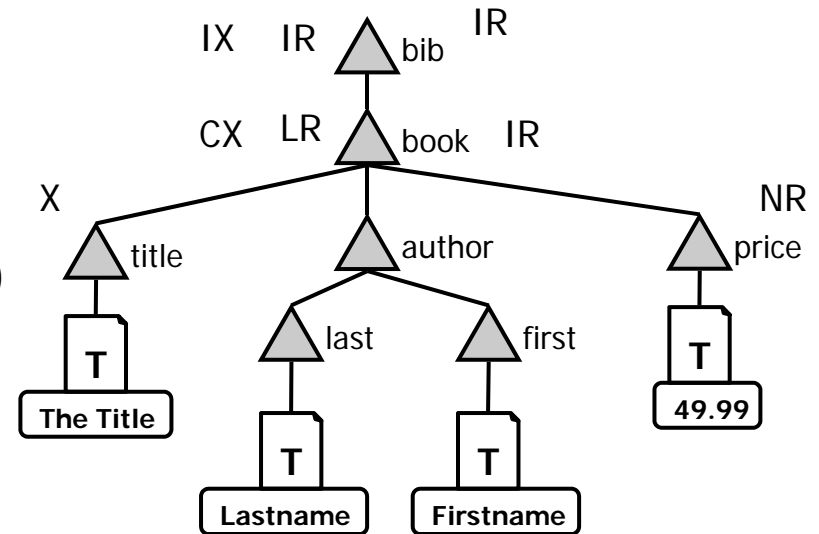
- 
- 
- 

# Node Locks

Transaction  $T_1$  is reading  $\langle \text{price} \rangle$

Transaction  $T_2$  is reading  $\langle \text{book} \rangle$   
and all direct-child nodes  
( $\langle \text{title} \rangle$ ,  $\langle \text{author} \rangle$ , and  $\langle \text{price} \rangle$ )

Transaction  $T_3$  is modifying  
the book title






- 
- 
- 

# Node Locks

- Locking subtrees exclusively: intention exclusive lock (IX), child exclusive lock (CX), and exclusive lock (X)
- Exclusive lock (X)
  - requested for updating the context node's content or deleting the context node and its entire subtree
  - requires a CX lock on the parent and IX locks on the ancestors



XML Transactions

- 
- 
- 
- 
- 
- 
- 
- 
-

- 
- 
- 

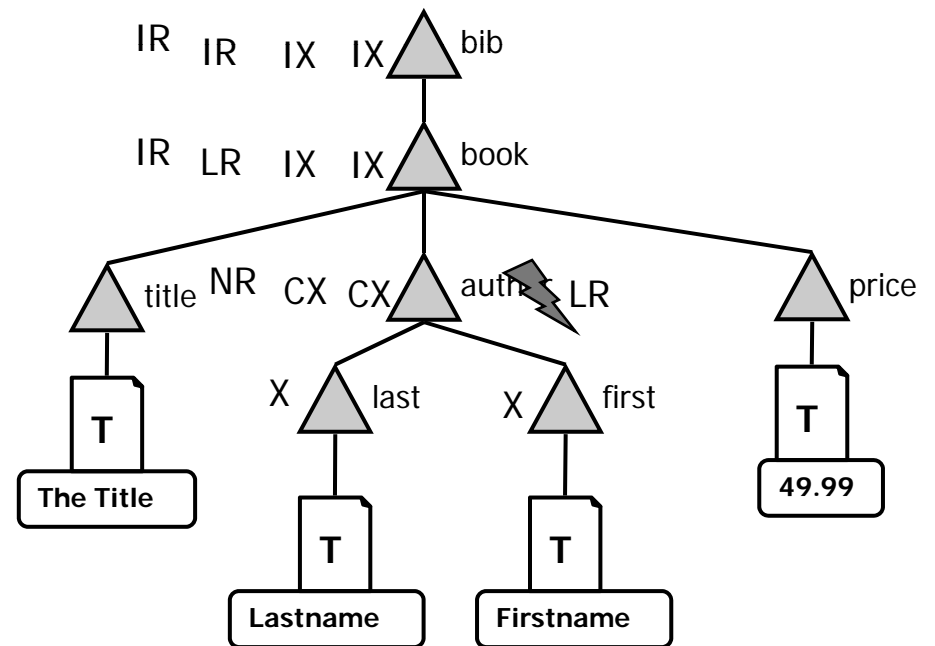
# Node Locks

Transaction  $T_1$  is deleting the `<first>` node and its content

Transaction  $T_2$  is deleting the `<last>` node and its content

Transaction  $T_3$  is reading the `<author>` node

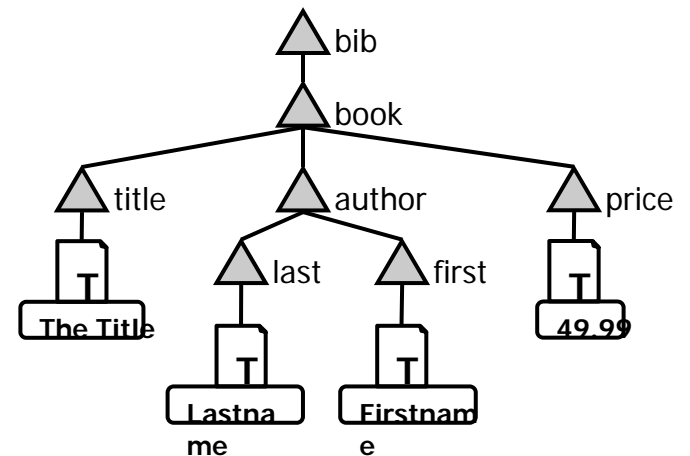
Transaction  $T_4$  is reading all direct-child nodes of `<author>` but is blocked when reading all child nodes of `<author>`



- 
- 
- 

# Tunable Lock Depth

- Goal
  - reduce the number of locks held by using coarser lock granularity
  - may decrease concurrency
  - when nodes deeper than lock depth are accessed: lock modes SR and X are used at the lock depth level



XML Transactions

35

• • • • • • • • • •

- 
- 
- 

# Tunable Lock Depth

would therefore have to wait on author

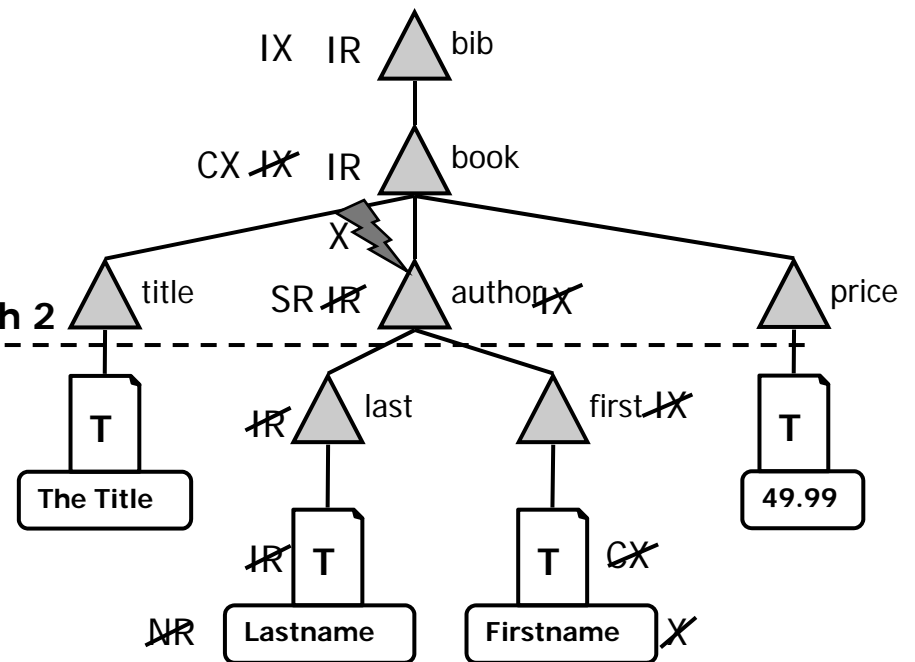
Transaction  $T_1$  is reading the author's last name

Transaction  $T_2$  is updating the author's first name

using lock depth 2

Transaction  $T_1$  would have to acquire an SR lock on author

Transaction  $T_2$  would have to acquire an X lock on author



XML Transactions

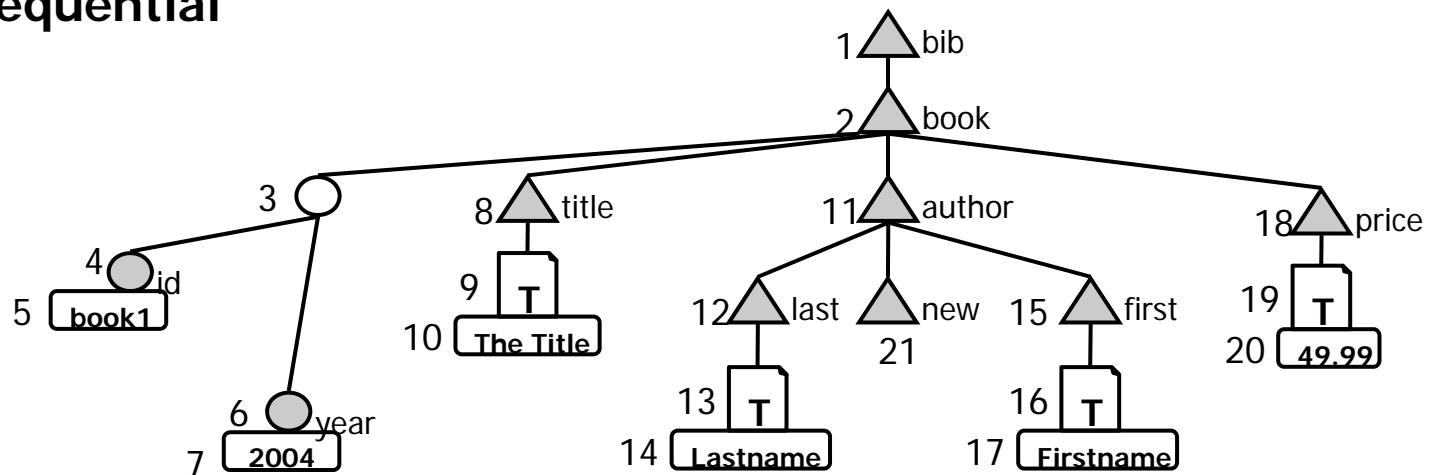
36

- 
- 
- 
- 
- 
- 
- 
-

•  
•  
•

## Identifying Nodes – Node Numbering Schemes

### Sequential

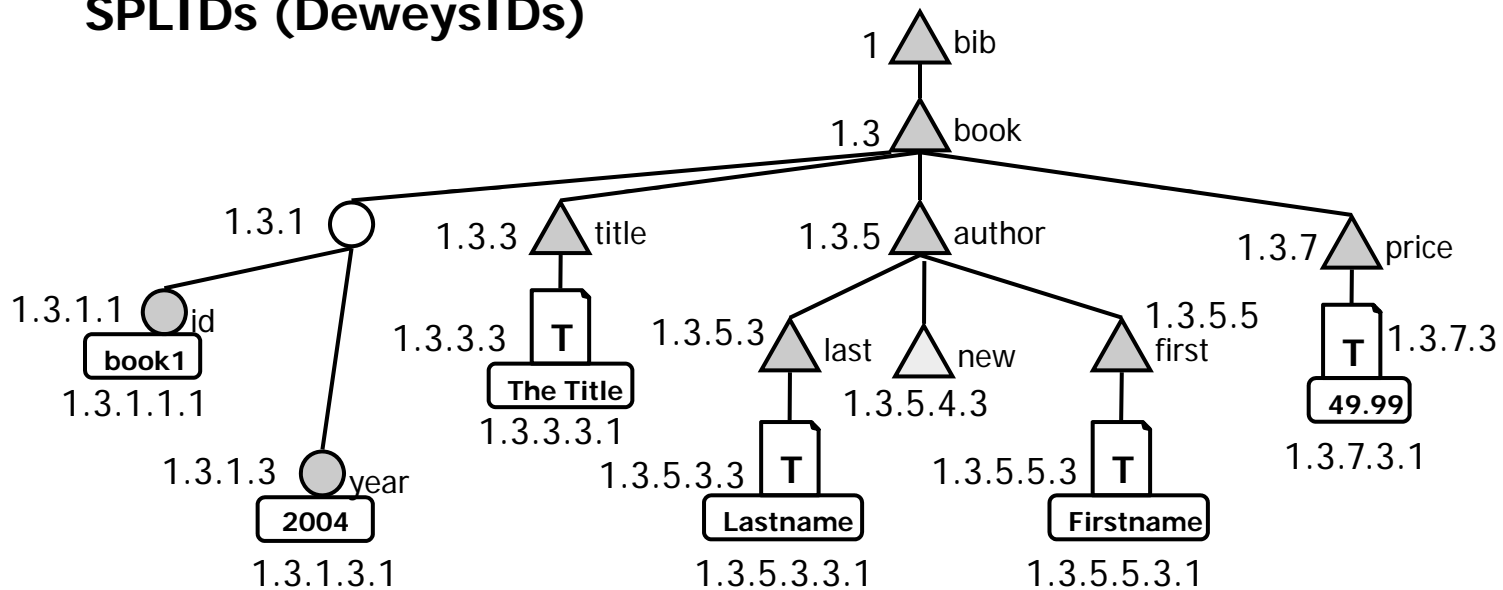


- very slow, although supported by on-demand indexes
- determination of parent ID and ancestor IDs, however, is very frequent

- 
- 
- 

## Identifying Nodes – Node Numbering Schemes

### SPLIDs (DeweysIDs)



- 
- 
- 

# OptiX/SnaX

- Locking overhead, especially for read operations, can be tremendous.
- Two snapshot based concurrency control mechanisms that avoid locking.
  - OptiX is a variation of optimistic concurrency control adjusted to use snapshots and work on XML data.
  - SnaX provides the isolation level of snapshot isolation and has similar semantics as the concurrency control mechanisms implemented
- Zeeshan Sardar, Bettina Kemme. “Don’t be a Pessimist: Use Snapshot based Concurrency Control for XML”, ICDE 2006.

•  
•  
•

# Snapshot based Concurrency Control

- Each transaction
  - Working phase, validation phase, update phase
- Two transactions  $T_i$ ,  $T_j$ 
  - Concurrent if  $T_i$  started the working phase before  $T_j$  finished the update phase and committed
- Validation
  - OptiX: WriteSet( $T_i$ ) not overlaps with ReadSet( $T_j$ )
  - SnaX: WriteSet( $T_i$ ) not overlaps with WriteSet( $T_j$ )
    - Snapshot isolation as in many RDBMS
    - Assuming that many applications are read only



•  
•  
•

## ReadSets

- XML structure
  - $RR(T_i)$  : roots of subtree returned as part of the query result
  - $ER(T_i)$  : nodes that are explicitly read as part of a predicate or path constraint, but not return as part of result
  - $ERa(T_i)$  : nodes that are read for the insertion of a sibling after them

•  
•  
•

## WriteSets

- XML structure
  - $D(T_i)$  : roots of subtree deleted or replaced
  - $R_n(T_i)$  : nodes to be renamed
  - $I(T_i)$  : the immediate parents of any nodes inserted by  $T_i$
  - $I_a(T_i)$  : the same set of nodes as in  $Era(T_i)$

- 
- 
- 

## OptiX Validation

- Suppose  $T_i$  wants to validate, and concurrent transaction  $T_j$  validated before  $T_i$ .
- When  $\text{ReadSet}(T_i)$  overlaps with  $\text{WriteSet}(T_j)$ 
  - Instead of immediately inducing a conflict, we look at the subsets of  $\text{ReadSet}(T_i)$  and  $\text{WriteSet}(T_j)$  of which  $p$  is a member.
  - YES indicates compatibility, while NO shows a conflict leading to an abort of  $T_i$ .
- A conflict means that if  $T_i$  had executed serially after  $T_j$ , then it would have possibly read a different value for  $p$  than in the concurrent execution.
- In the table,  $q$  is a descendant of  $p$ .

- 
- 
- 

# OptiX Validation

Tj already validated on p

Ti validating	D	Rn	I	Ia
On p RR	No	No	No	Yes
On p ER	No	No	Yes	Yes
On p ERa	No	No	Yes	No
On q RS	No	Yes	Yes	Yes

XML Transactions

44

- 
- 
- 

# SnaX Validation

- In SnaX, remember snapshot isolation does not guarantee that there is an equivalent serial execution.
  - ignore the read sets and only consider write/write conflicts.
- Generally, delete and replace conflict with most other update types, while inserts conflict only with few.

- 
- 
- 

# SnaX Validation

Tj already validated on p

Ti validating	D	Rn	I	la
On p D	No	No	No	Yes
On p Rn	No	No	Yes	Yes
On p I	No	Yes	Yes	Yes
On p la	Yes	Yes	Yes	No
On q WS	No	Yes	Yes	Yes