

•
•
•
•
•
•
•
•
•
•
•
•

XML Data Streams



•
•
•

Data Stream Processing

- What is a data stream?
 - continuous, time-varying data arriving at unpredictable rates
 - continuous updates, continuous queries
 - no stored index is available
- Sought characteristics of stream processing engines
 - real-time processing
 - high throughput, low latency, fast mean response time, low jitter
 - low memory footprint

-
-
-

Data Stream Processing

- Why bother?
 - many data are already available in stream form
 - sensor networks, network traffic monitoring, stock tickers
 - publisher-subscriber systems
 - data stream mining for fraud detection
 - data may be too volatile to index
 - continuous measurements


•
•
•

XML Stream Processing

- Various sources of XML streams
 - tokenized XML documents
 - sensor XML data
 - RSS feeds
 - web service results
 - MPEG-7 (binary encoding in XML)
- Granularity
 - XML tokens (events): `<tag>`, `</tag>`, “X”, etc
 - region-encoded XML elements
 - XML fragments (hole-filler model)

•
•
•

Traditional Stream Processing



- Typically, a stream consists of numerical values or relational tuples
- Focuses on a sliding window
 - fixed number of tuples, or
 - fixed time span
- Extracts approximate results
- Uses a small (bounded) state
- Examples:
 - top-k most frequent values
 - group-by SQL queries (OLAP)
 - data stream mining

• • • • • • • •

-
-
-

XML Update Streams

- A continuous (possibly infinite) sequence of XML tokens with embedded updates
 - Usually, a finite data stream followed by an infinite stream of updates
 - three basic types of tokens: `<tag>`, `</tag>`, `text`
 - the target of an update is a stream subsequence that contains zero, one, or more “complete” XML elements
 - the source is also a token sequence that contains complete XML elements

-
-
-

XML Update Streams

- the source is also a token sequence that contains complete XML elements
- updates are embedded in the data stream and can come at any time
 - update events can be interleaved with data events and with each other
 - each event must now have an id to associate it with an update
- updated regions can be updated too
- to update a stream subsequence, you wrap it in a Mutable region
- three types of updates:
 - replace
 - insertBefore
 - insertAfter

•
•
•

An Example

id	Event	... equivalent to
1	<a>	<a>
1		
1	StartMutable(2)	<c>
2	<c>	Y
2	X	</c>
2	</c>	<c>
1	EndMutable(2)	X
1		</c>
2	StartInsertBefore(3)	
3	<c>	
3	Y	
3	</c>	
2	EndInsertBefore(3)	
1		

• • • • • • • •

-
-
-

Continuous Queries

- Need to decide: snapshot or temporal stream processing?
 - Snapshot: after a replace update, the replaced element is forgotten
 - Temporal: “some” of the replaced elements are kept
 - we may have repeated updates on a mutable region, forming a history list
 - each version has a time span (valid begin/end times)
 - the versions kept are determined at run time from the temporal components of the query that process that region

-
-
-

Continuous Queries

- Query language: XQuery with temporal extensions
 - e?t time projection “give me the version before *t* secs”
 - e#v version projection “give me the past *v* version”
 - e?[t] time sliding window “give me all versions the last *t* secs”
 - e#[v] version sliding window “give me the *v* latest versions”
- The default is “current snapshot” (version #0 at time 0)
- Much finer grain for historical data than sliding windows

-
-
-

Continuous Results

- One can consider a stream engine is implemented as a pipeline
 - each pipeline stage performs a very simple task
- The final pipeline stage is the **Result Display** that displays the query results continuously
 - the display can be shown as a editable text window (a GUI), where text can be inserted, deleted, and replaced at any point
 - when an update is coming in the input stream, it is propagated through the result display, where it causes an update to the display text!

-
-
-

Snapshot Example

- *XQuery*

```
<books>{  
  for $b in stream("books")//biblio[publisher="Wiley"]/books  
  where $b/author/lastname="Smith"  
  order by $b/price  
  return <book>{ $b/title, $b/price }</book>  
}</books>
```

Display

```
<books>  
  <book><title>All about XML</title><price>35</price></book>  
  <book><title>XQuery for Dummies</title><price>58</price></book>  
  <book><title>Querying XML</title><price>120</price></book>  
  ...
```

-
-
-

A Temporal Query

- Display all stocks whose quotation increased at least 10% since the last time, sorted by their rate of change:

```
<quotes>{  
    for $q in stream("tickers")//ticker  
    where $q/quote > $q/quote#1 * 1.1  
    order by ($q/quote - $q/quote#1) div $q/quote  
    return <quote>{ $q/name, $q/quote }</quote>  
}</quotes>
```

-
-
-

Efficient Evaluation of XQuery over Streaming Data

- XPath over Streaming Data
 - XPath is relatively simple
- XQuery over Streaming Data
 - Limited features handled
 - Focus on queries that are written for single pass evaluation
- VLDB 2005
 - Xiaogang Li , Gagan Agrawal

-
-
-

Ideas

- Can the given query be evaluated correctly on streaming data?
 - Only a single pass is allowed
 - Decision made by compiler, not a user
- If not, can it be correctly transformed ?
- How to generate efficient code for XQuery?
 - Computations involved in streaming application are non-trivial
 - Recursive functions are frequently used
 - Efficient memory usage is important

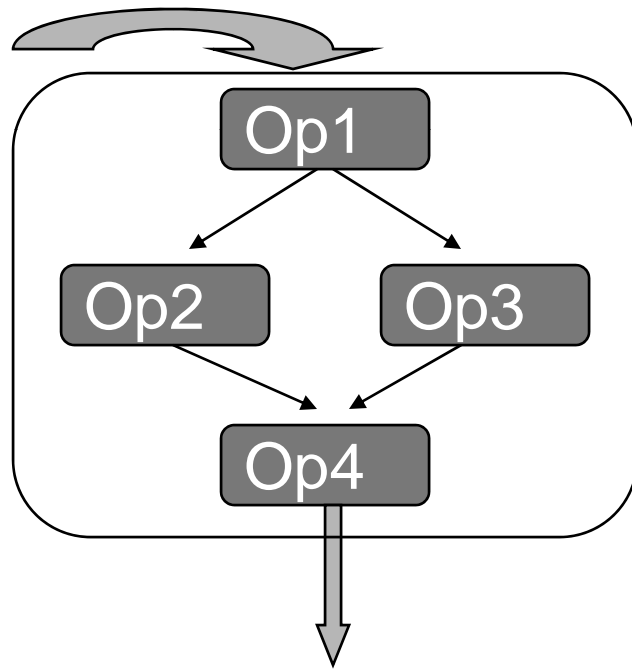
-
-
-

The Approach

- For an arbitrary query, can it be evaluated correctly on streaming data?
 - Construct data-flow graph for a query
 - Static analysis based on data-flow graph
- If not, can it be transformed to do so ?
 - Query transformation techniques based on static analysis
- How to generate efficient code for XQuery?
 - Techniques based on static analysis to minimize memory usage and optimize code
 - Generating imperative code
 - Recursive analysis and aggregation rewrite

•
•
•

Query Evaluation Model



- Single input stream
- Internal computations
 - Limited memory
 - Linked operators
- Pipeline operator and Blocking operator

-
-
-

Pipeline and Blocking Operators

- Pipeline Operator:
 - each input tuple produces an output tuple independently
 - Selection, Increment etc
- Blocking Operator:
 - Can only compute output after receiving all input tuples
 - Sort, Join etc
- Progressive Blocking Operator:
 - (1) $|output| \ll |input|$: we can buffer the output
 - (2) Associative and commutative operation: discard input
 - count(), sum()

•
•
•

Single Pass?

Pixels with x and y

Q1:

let \$i := .../pixel

sortby (x)

(1) A blocking operator exists

Q2:

let \$i := for \$p in /pixel

where \$p/x > ..

x = count(/pixel)

(2) A progressive blocking operator is referred by another pipeline operator or progressive operator

Check condition 2 in a query

• • • • • • • •

-
-
-

Challenges in Single-Pass

```
let $b = count(stream/pixel[x>0])  
  for $i in stream/pixel  
    return $i/x idvi $b
```

Must Analyze data dependence at expression level

```
let $b: = for $i in stream/pixel[x>0]  
  return $i  
for $j in $b/y  
  return $j  
  where $j = count($b)
```

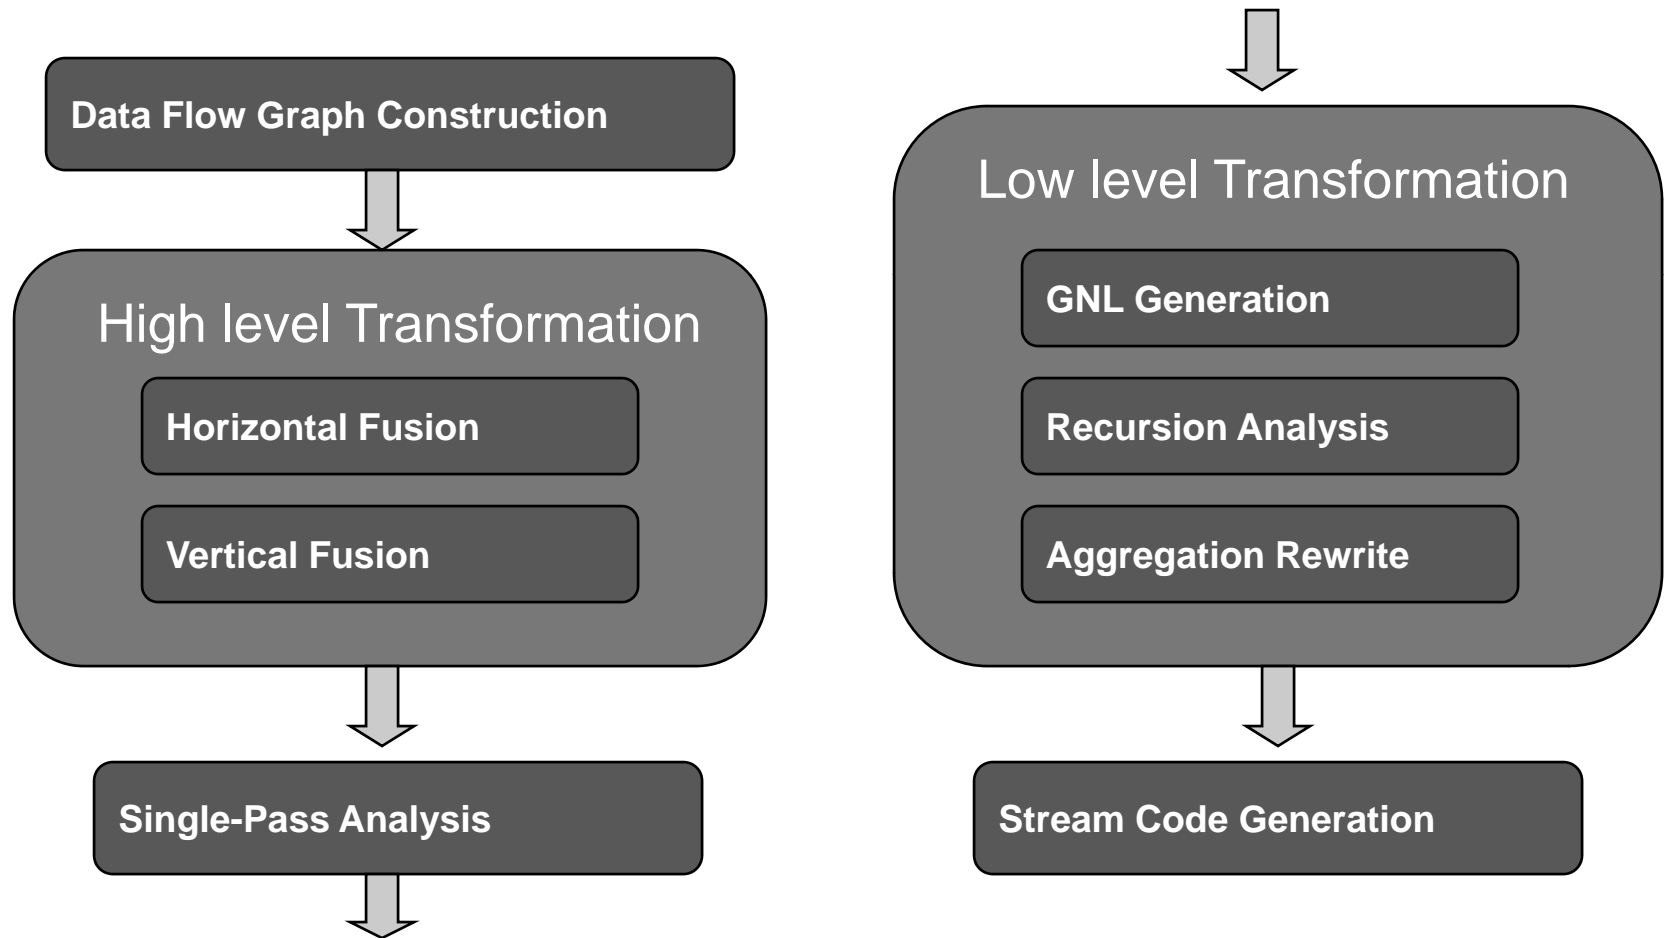
A Query may be complex

Need a simplified view of the query to make decision

•

•


•



•
•
•

Stream Data Flow Graph (DFG)

```
let $b = count(stream/pixel[x>0])  
for $i in stream/pixel  
  return $i/x idvi $b
```



S1:stream/pixel[x>0]

S2:stream/pixel

V1: count()

- Node represents variable:
Explicit and implicit
Sequence and atomic

- Edge: dependence relation
 $v1 \rightarrow v2$ if $v2$ uses $v1$
Aggregate dependence and
flow dependence

- A DFG is acyclic
- Cardinality inference is
required to construct the
DFG

• • • • • • • •

-
-
-

High-level Transformation

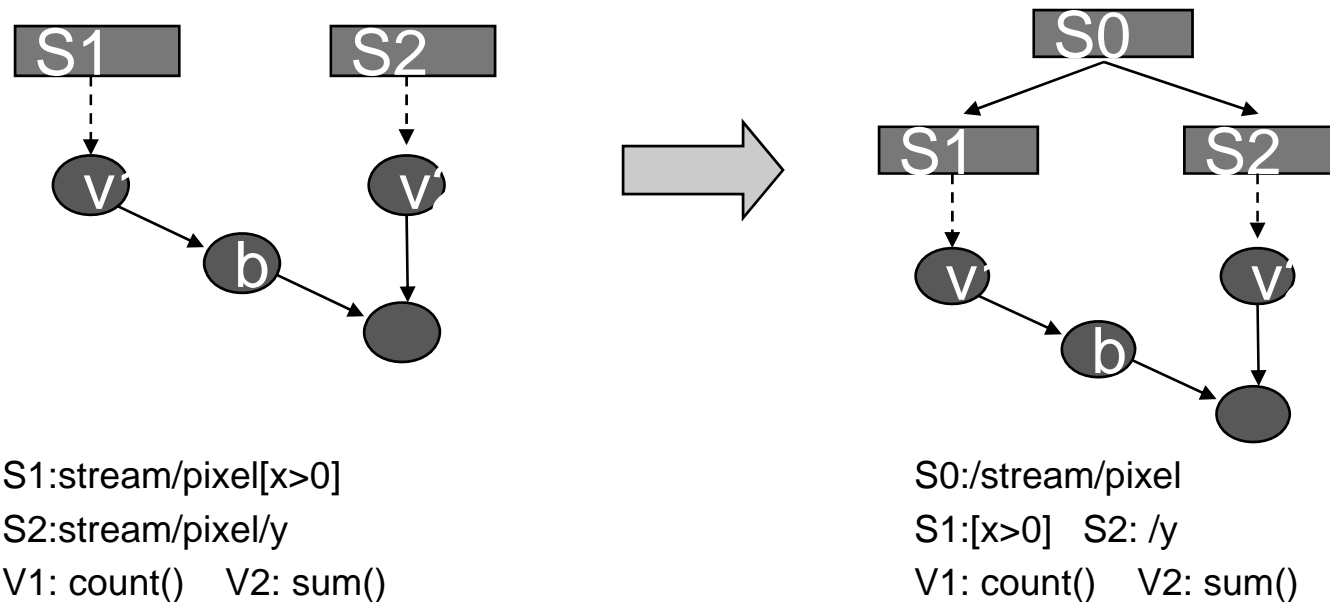
- Goals
 - 1: Enable single pass evaluation
 - 2: Simplify the SDFG and single-pass analysis
- Horizontal Fusion and Vertical Fusion
 - Based on SDFG

•
•
•

Horizontal Fusion

- Enable single-pass evaluation
 - Merge sequence node with common prefix

```
let $b = count(stream/pixel[x>0])
return sum(stream/pixel/y) idvi $b
```




• • • • • • • •

-
-
-

Horizontal Fusion with Nested Loops

- Perform loop unrolling first
- Merge sequence node accordingly

```
unordered(  
  for $i in (1 to 2)  
    let $b: //=stream/pixel[x=$i]  
    return count($b))
```




```
unordered(  
  let $b1: //=stream/pixel[x=1]  
  let $b2: //=stream/pixel[x=2]  
  return count($b1), count($b2))
```

-
-
-

Horizontal Fusion: Side-effect

- May result in incorrect result due to inter-dependence




Partial result of count is used to compute output
Will be dealt with at single-pass analysis

•
•
•

Vertical Fusion

- Simplify DFG and single-pass analysis
 - Merge a cluster of nodes linked by flow dependence edges

```
let $b: = for $i in stream/pixel[x>0]
  return $i
for $j in $b/y
  return $j
where $j = count($b)
```



• • • • •

•
•
•

Single-pass Analysis

- Can a query be evaluated on-the fly?

THEOREM 1. *If a query with dependence graph $G=(V,E)$ contains more than one sequence node after vertical fusion, it can not be evaluated correctly in a single pass.*


Reason: Sequence node with infinite length can not be buffered

•
•
•

Single-pass Analysis

THEOREM 2. *Let S be the set of atomic nodes that are aggregate dependent on any sequence node in a stream data flow graph. For any given two elements $s1$ and $s2$, if there is a path between $s1$ and $s2$, the query may not be evaluated correctly in a single pass.*

Reason: A progressive blocking operator is referred by another progressive blocking operator



Example : count (pixel)
where $/x > 0.005 * \text{sum}(/ \text{pixel} / x)$


• • • • • • • •

•
•
•

Single-pass Analysis

THEOREM 3. *In there is a cycle in a stream data flow graph G , the corresponding query may not be evaluated correctly using a single pass.*

Reason: A progressive blocking operator is referred by a pipeline operator



• • • • •

-
-
-

Single-pass Analysis

- Check conditions corresponding to Theorem 1 2 and 3
 - Stop further processing if any condition is true
- Completeness of the analysis
 - If a query without blocking operator pass the test, it can be evaluated in a single pass

THEOREM 4. *If the results of a progressive blocking operator with an unbounded input are referred to by a pipeline operator or a progressive blocking operator with unbounded input, then for the stream data flow graph, at least one of the three conditions holds true*

•
•
•

Conservative Analysis

- Our analysis is conservative
 - A valid query may be labeled as “cannot be evaluated in a single-pass”

Example:

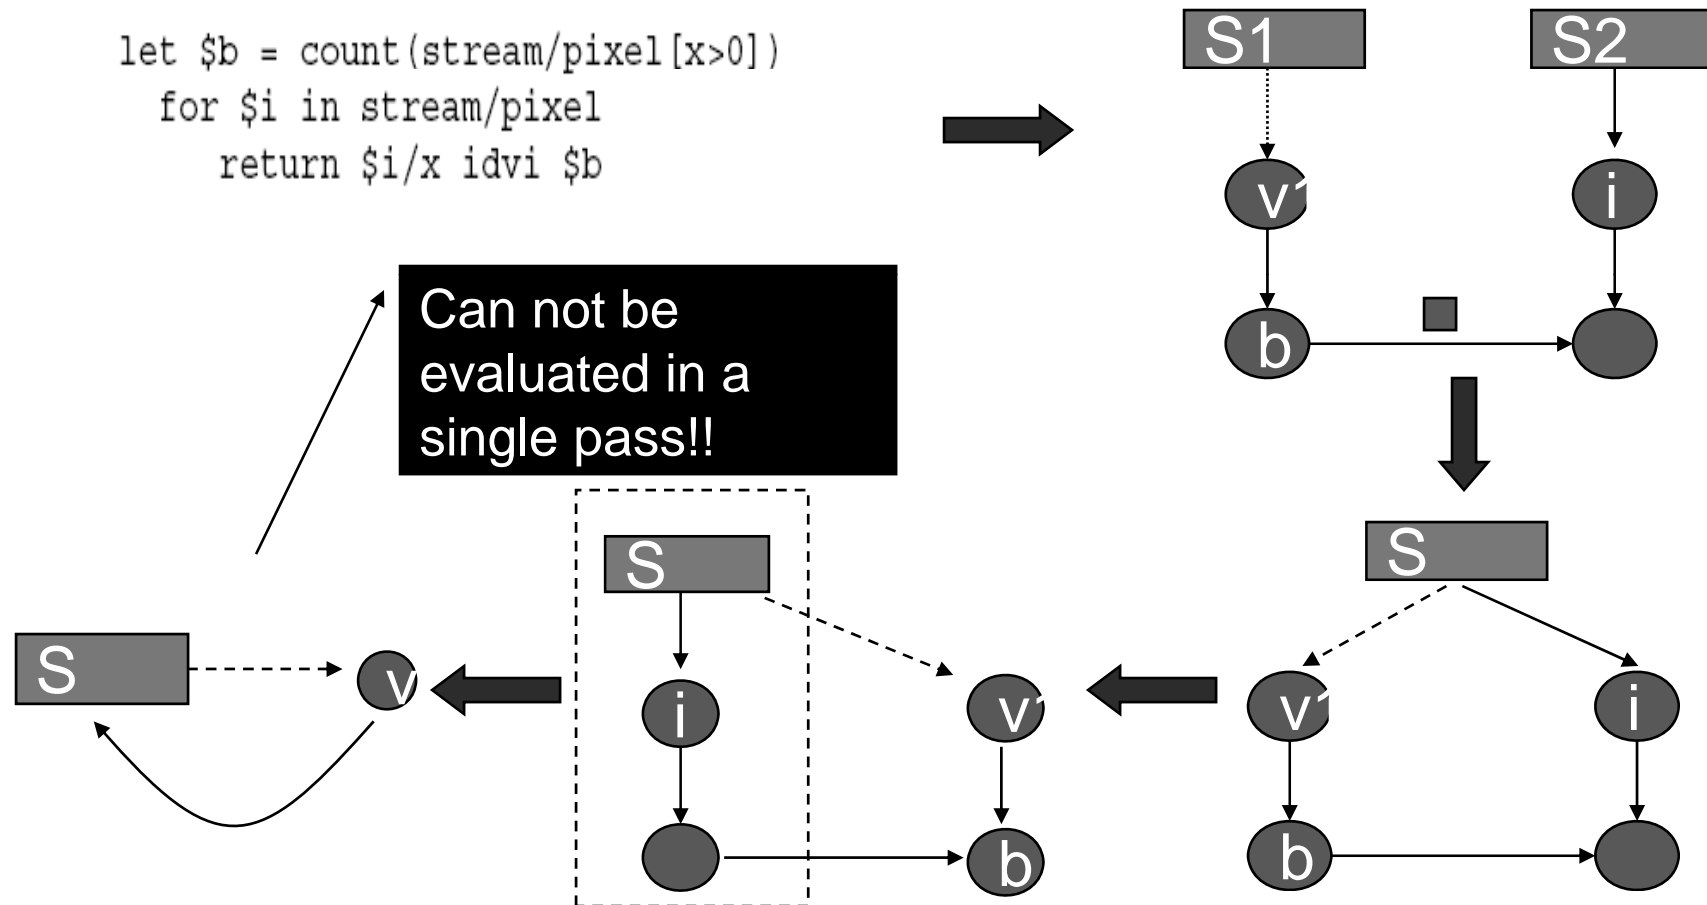
```
let $p: = stream/pixel/x
  for $i in $p
    where $i <= max($p)
  return $i
```

• • • • • • • •

•
•
•

The Overall Procedure

```
let $b = count(stream/pixel[x>0])  
for $i in stream/pixel  
  return $i/x idv $b
```



-
-
-


Low-level Transformations

- Use GNL as intermediate representation
 - GNL is similar to nested loops in Java
 - Enable efficient code generation for reductions
 - Enable transformation of recursive function into iterative operation
- From SDFG to GNL
 - Generate a GNL for each sequence node associated with XPath expression
 - Move aggregation into GNL using aggregation rewrite and recursion analysis

•
•
•

GNL Example

```
let $b = count(stream/pixel[x>0])
return sum(stream/pixel/y) idvi $b
```



```
for i1, stream/pixel, --
[ for i2, /x, /x > 0
  [ v1 = v1 + 1
    for i3, /y, --
      [ v2 = v2 + i3 ;
        b = v1 return b ÷ v2
```

Facilitate code generation for any desired platform

-
-
-

Low-Level Transformations

- Recursive Analysis
 - extract commutative and associative operations from recursive functions
- Aggregation Rewrite
 - perform function inlining
 - transform built-in and user-defined aggregate into iterative operations

-
-
-

Code Generation

- Using SAX XML stream parser
 - XML document is parsed as stream of events
`<x> 5 </x>`: startelement `<x>`, content 5, endelement `<x>`
 - Event-Driven: Need to generate code to handle each event
- Using Java JDK
 - Our compiler generates Java source code

-
-
-

Code Generation: Example

```

for  $i_1$ , stream/pixel, --
[ for  $i_2$ , /x, /x > 0
  [  $v_1 = v_1 + 1$ 
    for  $i_3$ , /y, --
      [  $v_2 = v_2 + i_3$  ;
     $b = v_1$  return  $b \div v_2$ 
  ]
]

```

startElement: Insert each
referred element into buffer

endElement: Process each
element in the buffer,
dispatch the buffer

```

foreach startElement ( $e_i$ ) {
  switch(  $e_i$ .node)
    x: buffer.add(x)
    y: buffer.add(y)
}

foreach endElement ( $e_i$ ) {
  switch(  $e_i$ .node)
    x: if (buffer.dispatch(x) > 0)
       $v_1 = v_1 + 1$ 
    y:  $v_2 = v_2 +$  buffer.dispatch(y)
    root: {  $b = v_1$ ;
      return  $b / v_2$ 
    }
}

```