

-
-
-
-
-
-
-
-
-
-

XML Transactions



-
-
-

XML as Database

- Multiple Users
 - 1~1000, or more?
 - Querying and updating occur simultaneously
- Transaction Management
 - Atomicity of query and update operations
 - All-or-nothing execution
 - Consistency and Concurrency Control
 - Locking system

-
-
-

XQuery Update

- XQuery Update Facility 1.0
 - <http://www.w3.org/TR/2008/CR-xquery-update-10-20080801/>
- Primitive Updates: insert, delete, replace, rename
 - Extensions to other expr: FLWOR, TypeSwitch, ...

-
-
-

Examples

- delete nodes `//book[@year lt 1968]`
- insert node `<author/>` into `//book[@ISBN eq "34556"]`
- for `$x` in `//book`
 where `$x/year lt 2000` and `$x/price gt 100`
 return replace value of node `$x/price`
 with `$x/price-0.3*$x/price`
- if (`$book/price gt 200`) then
 rename node `$book` as "expensive-book"

Update statements work on "node" or "nodes"

-
-
-

Insert

- Insert a new element into a document
- **insert node UpdateContent into TargetNode**
 - UpdateContent: any sequence of items (nodes, values)
 - TargetNode: Exactly one document or element
 - Otherwise ERROR
- Optionally, specify if to insert at the beginning or end
 - as last: Content becomes first child of Target
 - as first: Content becomes last child of Target
 - No position: no fixed position (honor other first/last inserts)
- Nodes in Content assume a new Id.

-
-
-

Insert

- Insert new book in the library
insert node <book> <title>Die wilde Wutz</title> </book>
into document("www.uni-bib.de")//bib
- Insert new book at the beginning of the library
insert node <book> <title>Die wilde Wutz</title> </book>
as first into document("www.uni-bib.de")//bib
- Insert new attribute into an element
insert node (attribute age { 13 })
into document("persons.xml")//person[@name = "KD"]

-
-
-

Insert

- Insert at a particular point in the document
- **insert node UpdateContent (after | before) TargetNode**
 - UpdateContent: No attributes allowed!
 - TargetNode: One Element, Comment or PI.
 - Otherwise ERROR
 - Must have parent
- Specify whether before or behind target
 - Before vs. After
- Nodes in Content assume new identity

Add an author to a book

```
insert node <author author>Florescu Florescu</ </author author>  
before // article[title = "XL"]/ [author . = " Vincent"]
```

-
-
-

Delete

- Deletes nodes from a document
- **delete (node | nodes) TargetNodes**
 - TargetNodes: Any sequence of nodes
- Delete papers.
delete node //article[header/keyword = "XML"]

-
-
-

Replace

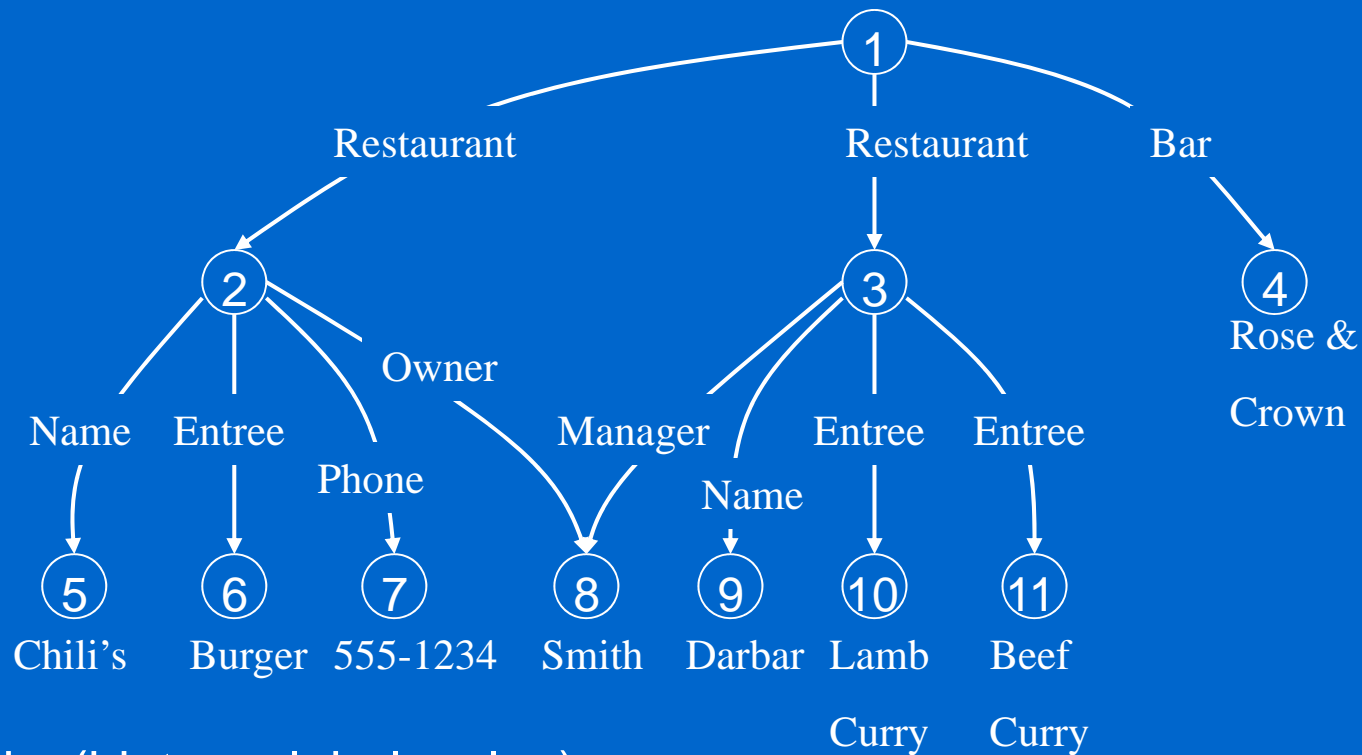
- Replace a node
replace node TargetNode **with** UpdateContent
- Replace the content of a node
replace value of node TargetNode **with** UpdateContent
 - TargetNode: One node (with Id)
 - UpdateContent: Any sequence of items
 - Keeps the node ID of TargetNode

-
-
-

Concurrency Control

- Concurrency control is a method used to ensure that database transactions are executed in a safe manner.
- ACID properties
 - **A** tomicity: All actions in the transaction happen, or none happen.
 - **C** onsistency: If each transaction is consistent, and the DB starts consistent, it ends up consistent.
 - **I** solation: Execution of one transaction is isolated from that (the effect) of other transactions.
 - **D** urability: If a transaction commits, its effects persist.


Dataguide model




nodes(id, type, label, value)

child(child-id, parent-id)

Dataguide model



Data-value Conflict



Access Path $L=/a/b/c$

T1: R(L) $r(a)r(b)r(c)$


T2: W(L) $r(a)r(b)w(c)$

$r(a)r(b) r(a)r(b) w(c) r(c)$ abort

Structural Conflict

L1=/a/d/e

L2=/a/d




T1: Append(L1,v) - $r(a)r(d)r(e)w(v)w(e)$

T2: Remove(L2) - $r(a)r(d)w(d)$

$r(a) r(a) r(d)r(e)w(v) r(d)w(d) w(e)$

Element-order Conflict



L-/a/g/f

T1: Append(L, u) $r(a)r(g)r(f)w(u)w(f)$

T2: Append(L, v) $r(a)r(g)r(f)w(v)w(f)$

What is a possible problem?


-
-
-

Data-value Conflict

L1=/a/c/d


L2=/a/b/d

index



T1: Read(L1)

T2: Write(L2)



Problem in the index
tree, common D node

-
-
-

Structural Conflict and Index Node Split



L1=/a/b/d

L2=/a/b

T1: Append(L1, v)

T2: Remove(L2) Rename(F)

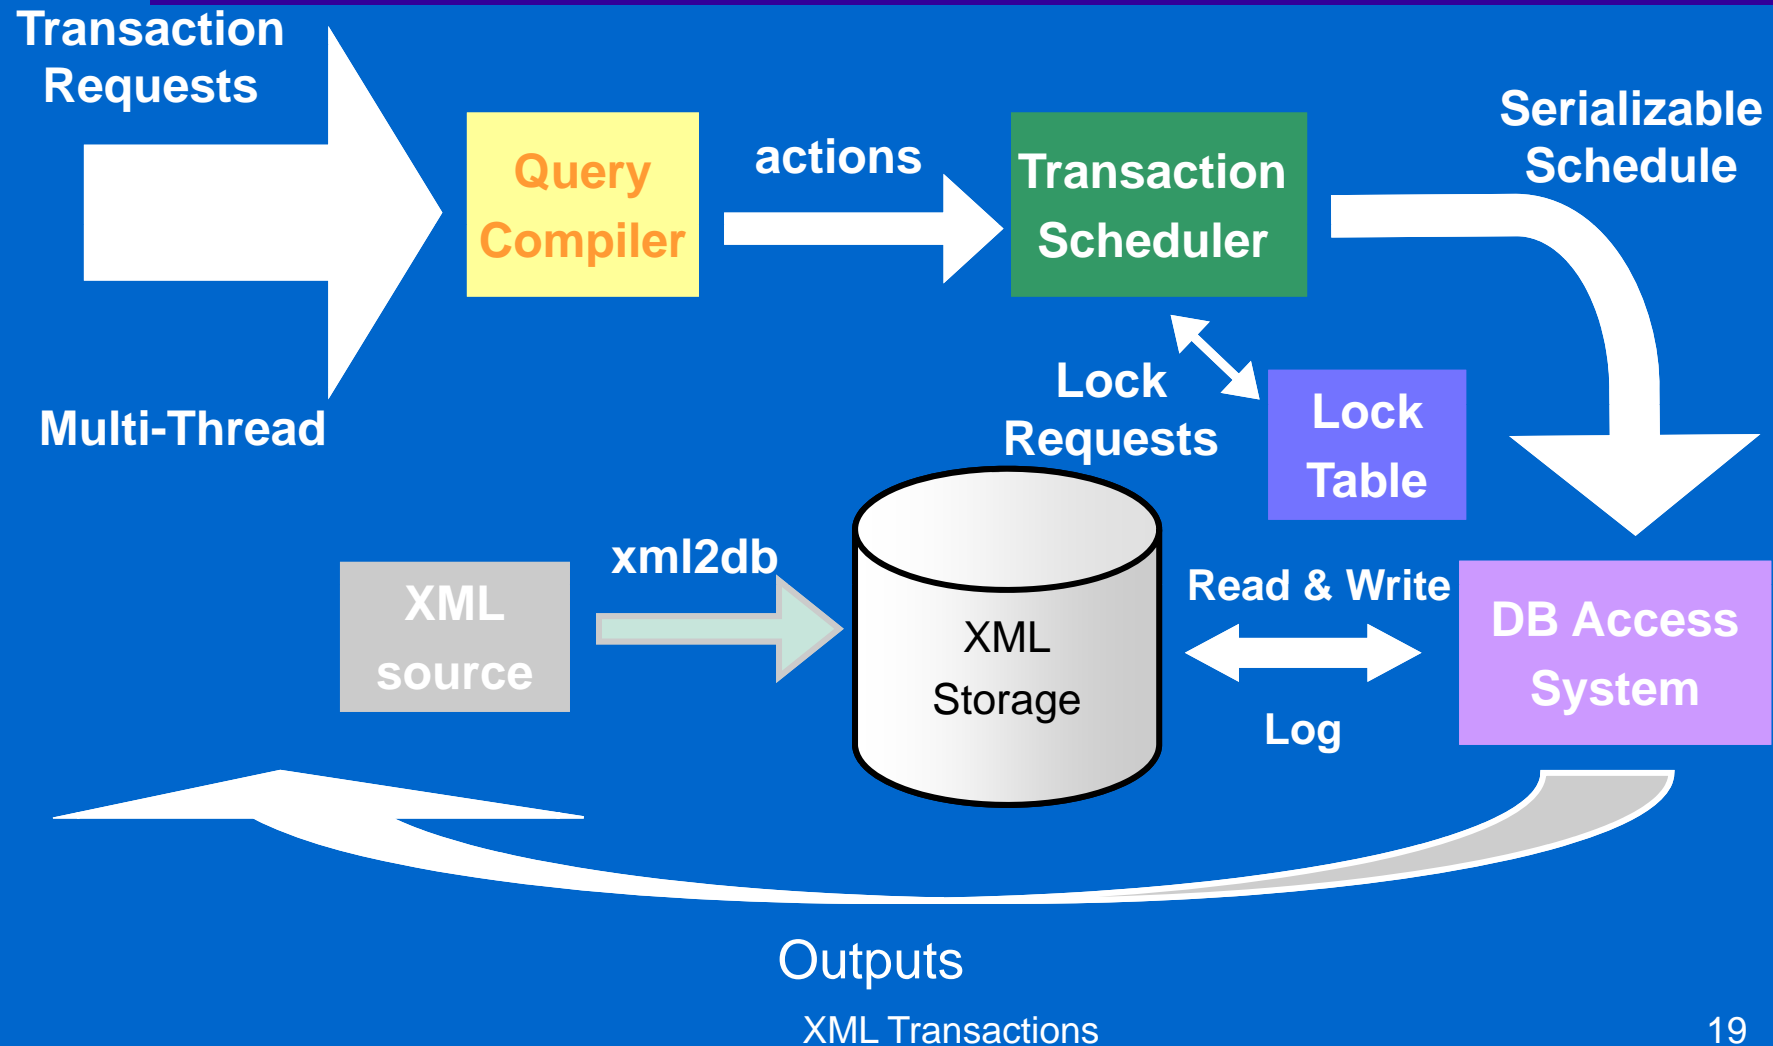
index



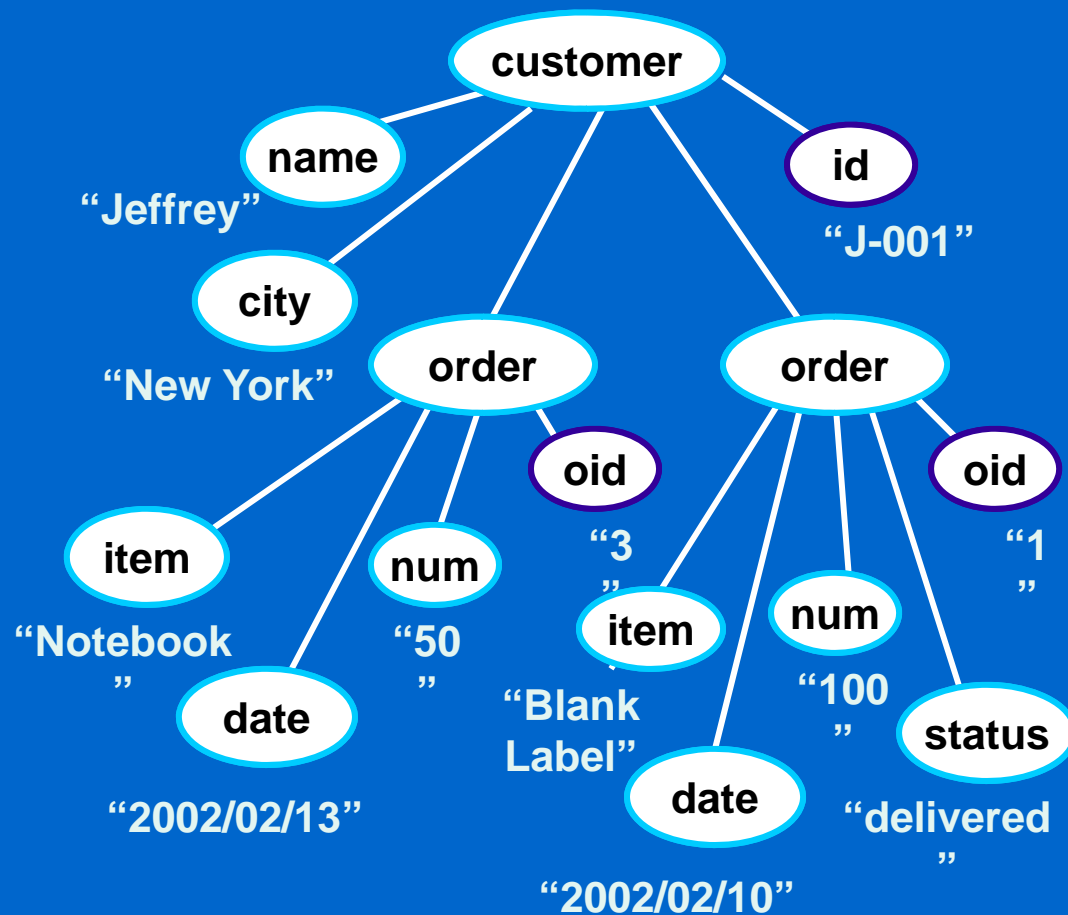
Xerial

- www.xerial.org
 - The name comes after XML and serializable
 - Taro L. Saito. On concurrency and updatability of XML databases. Master thesis submitted to Department of Computer Science, Graduate School of Information Science and Technology, University of Tokyo, January 2004
- Transactional Database for XML
 - Concurrent Transactions
 - Serializable schedule
 - Recoverability
 - Handling transaction aborts and system failures
 - Updating XML
 - Node insertion, deletion, modification, etc.
 - Transaction Language
 - Query and update notations

Xerial Overview



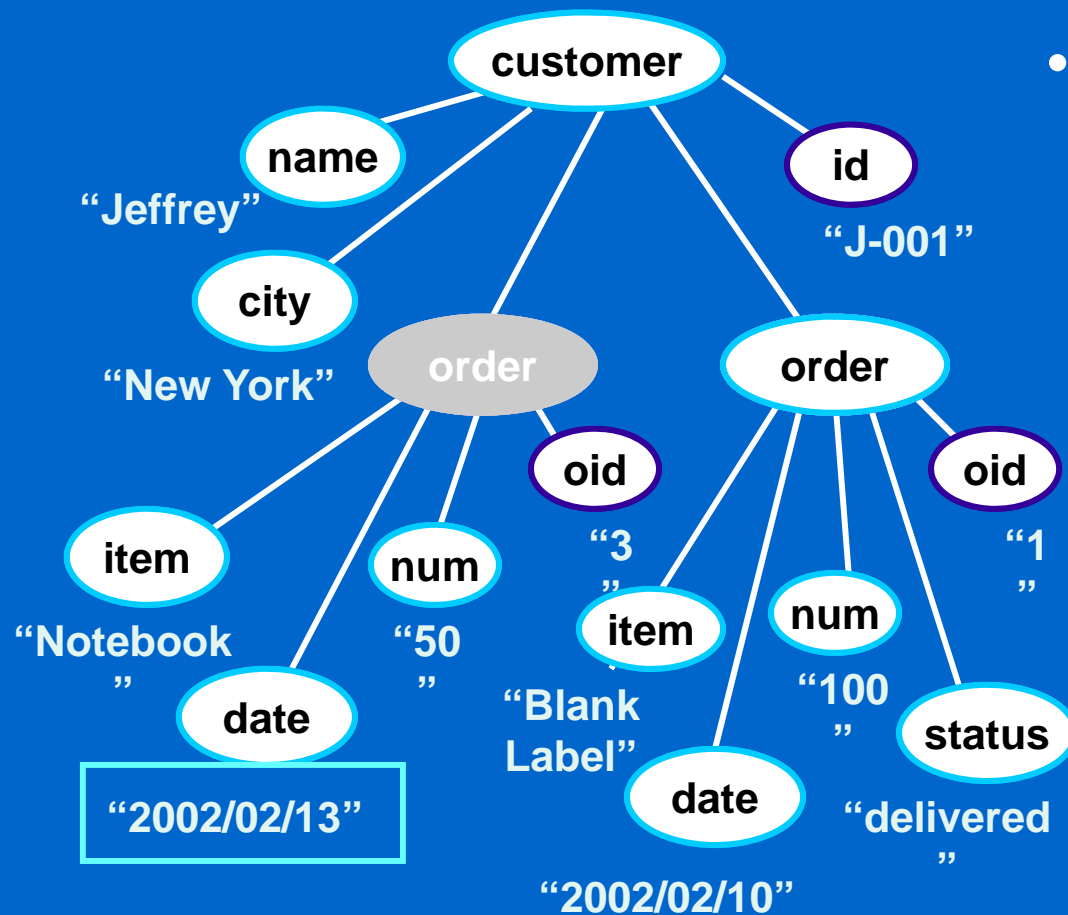
Data Model



```

<customer id="J-001">
  <name> Jeffrey </name>
  <city> New York </city>
  <order oid="3">
    <item> Notebook </item>
    <date> 2002/02/11 </date>
    <num> 50 </num>
  </order>
  <order oid="1">
    <item> Blank Label </item>
    <date> 2002/02/10 </date>
    <num> 100 </num>
    <status> delivered </status>
  </order>
</customer>
  
```

Querying XML




- **XQuery**

- W3C standard
- Query Language for XML
- Use of **Path expressions**
- Bind elements to a **variable**

FOR \$x **IN** /customer/order

WHERE \$x/date = "2002/02/13"

Locks for Tree-Structure



- **Subtree Level Locking**
 - Query to entire subtree is frequent in XML
 - Reduce the number of locks
- **Performance Factor**
 - The number of locks
 - Load of lock manager
 - Granularity of locks
 - Concurrency

-
-
-

Locks

t1/t2	IS	IX	S	X
IS	Yes	Yes	Yes	No
IX	Yes	Yes	No	No
S	Yes	No	Yes	No
X	No	No	No	No

Compatibility Matrix


Ordinal Locks

- *S* Shared Lock (read)
- *X* Exclusive Lock (write)

Warnings


- *IS* Intention to Share
- *IX* Intention to Exclusive

Warning Protocol



- Jim Gray et al, 1975.
- Original Rules
 - All transactions must enter from the root
 - To place a lock or warning on any element, we must hold a *warning* on its parent
 - Never remove a lock or warning unless we hold no locks or warnings on its children

Warning Protocol for XML



- **Extension**

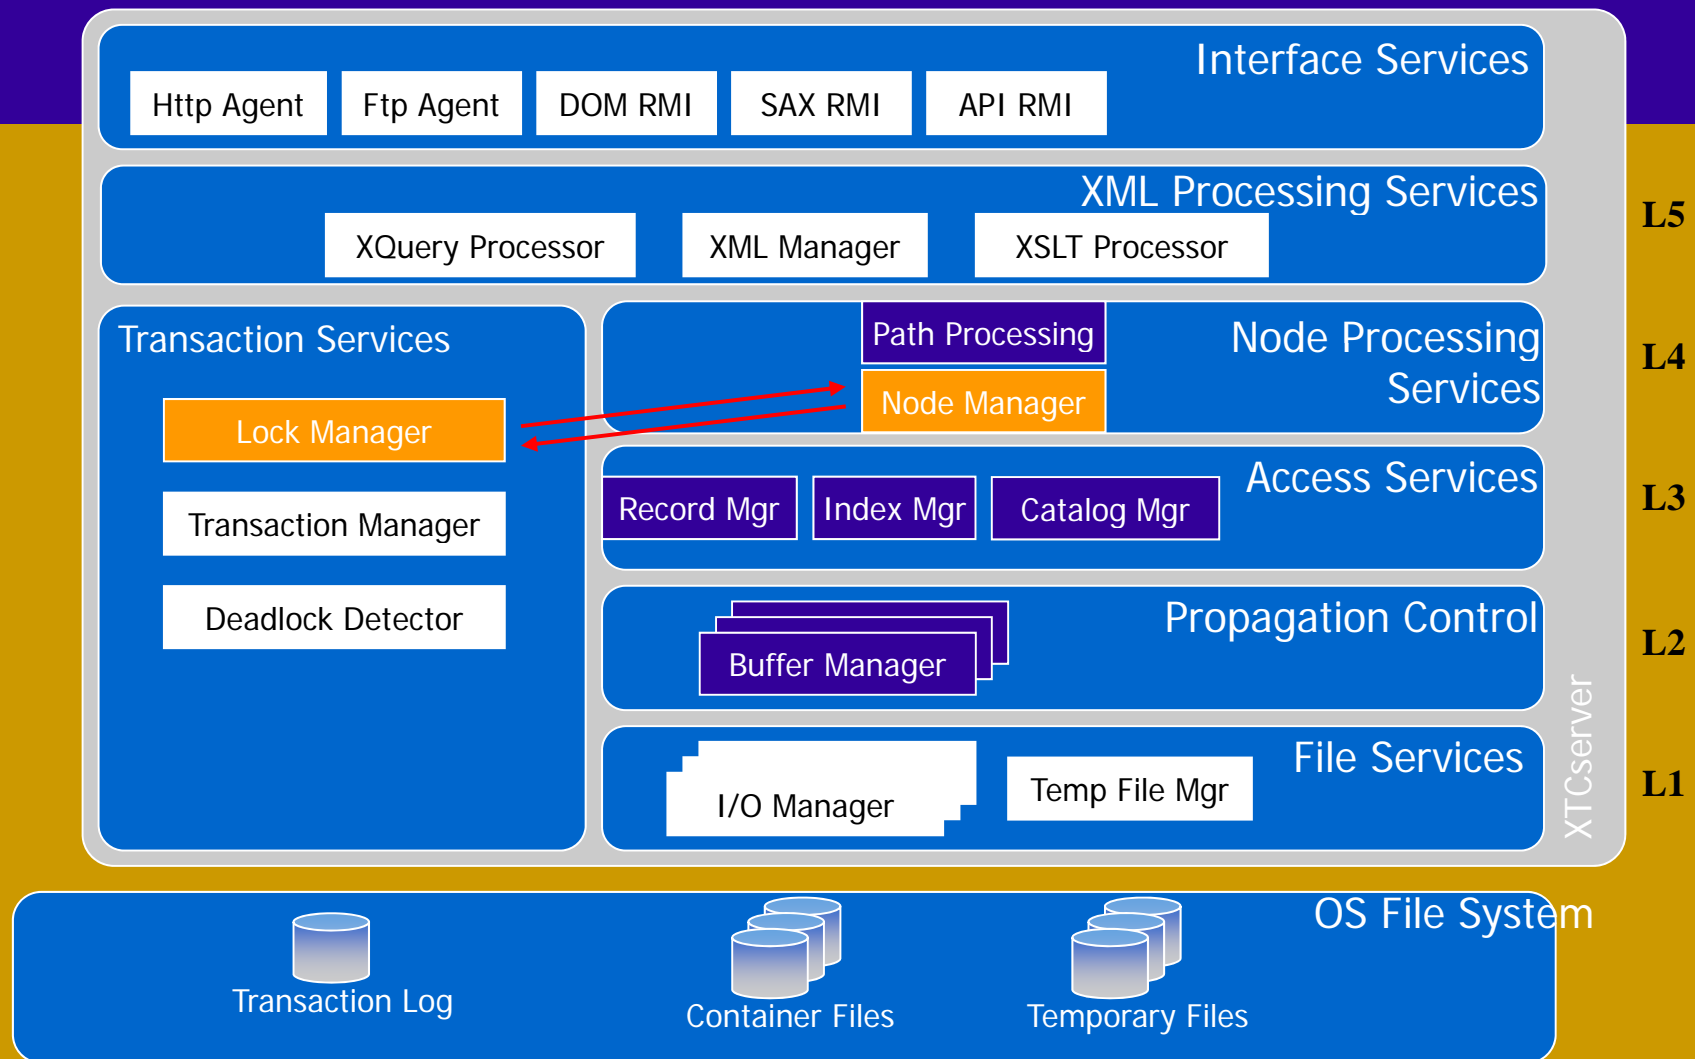
- When we insert or delete nodes, we must obtain *X* lock on the *parent* of the destination
- Until we place a warning on a node, we cannot trace its pointers to the children
- A transaction never releases locks or warnings until it finishes
 - **2 phase locking**

-
-
-

XTC

- XTC is used as a test vehicle for empirical DB research
 - effectiveness of XML concurrency control
 - fine-granular locking on nodes and edges
 - meta-synchronization allows comparison of different compatibilities
 - taDOM* protocols
 - multiplicity of lock modes
 - intention locks are important
 - indexed document access is frequent
 - ancestor path locking without accessing the storage engine desirable
- References
 - Michael P. Haustein, Theo Haerder: Twig Query Processing Under Concurrent Updates, in ICDE. 2006.
 - Michael P. Haustein, Contest of XML Lock Protocols, VLDB 2006 (www.vldb.org/conf/2006/p1069-haustein.pdf)

XTC – Architectural Overview



determination of ancestor node IDs are of outmost importance
for any locking protocol XML Transactions

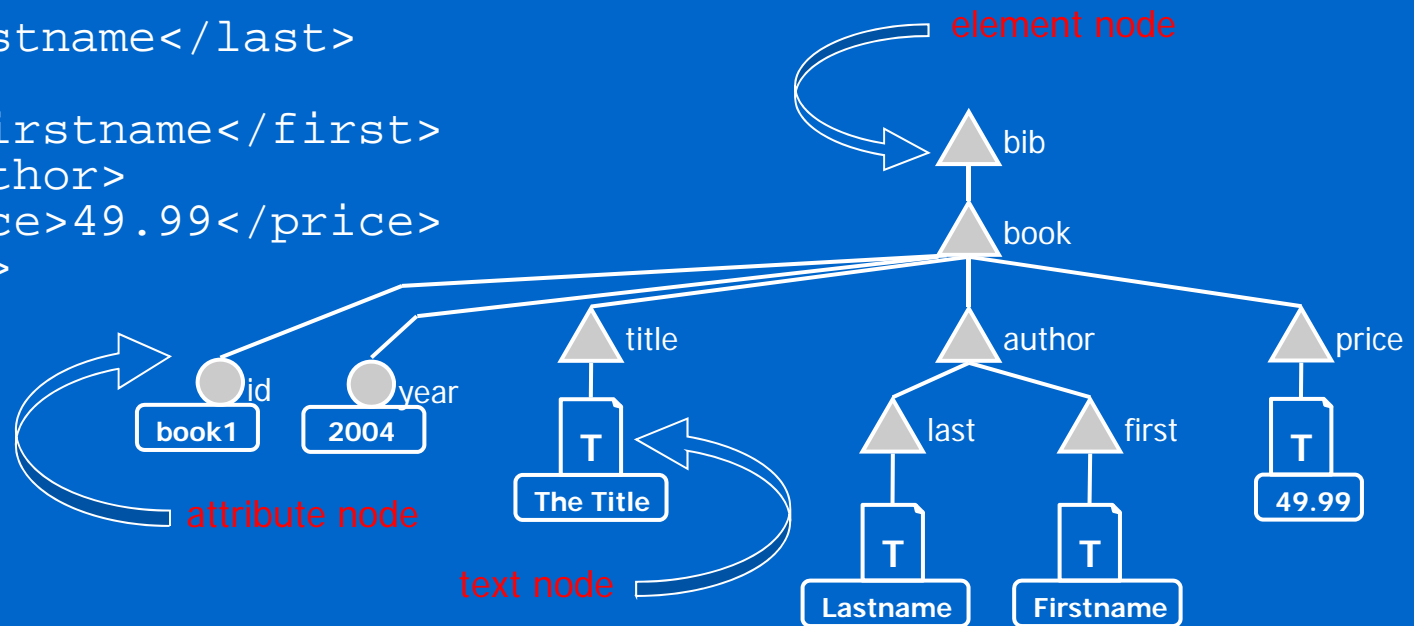
-
-
-

XML Document

```
<?xml version="1.0"?>
<bib>
  <book year="2004"
id="book1">
    <title>The
Title</title>
    <author>

<last>Lastname</last>

<first>Firstname</first>
    </author>
    <price>49.99</price>
  </book>
</bib>
```



Tailored Node Locks for XML – taDOM2

Read locks

lock	effect
IR	intention read lock on a node
NR (node read)	locks only a context node
LR (level read)	read lock on a context node and all direct-child nodes
SR (subtree read)	read lock on an entire subtree

Write locks

lock	effect
IX (intent. exclusive)	intention of a write lock on a non-direct child node
X (exclusive)	write lock on an entire subtree
CX (child exclusive)	indicates existence of a write lock on a direct child node
SU (update option)	read lock for intended update operations on an entire subtree

-
-
-

Tailored Node Locks for XML – taDOM2

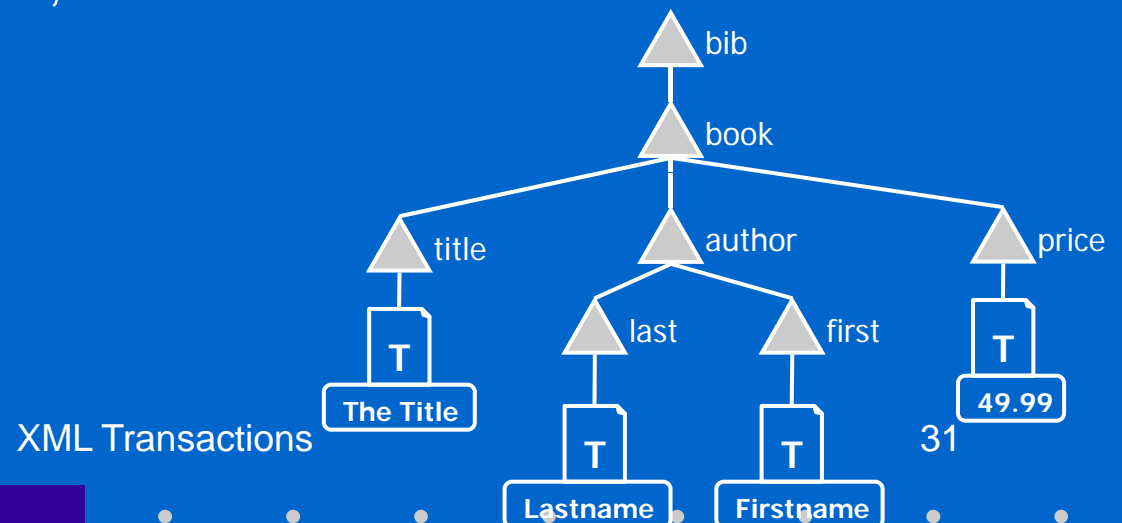
Compatibility matrix

	-	IR	NR	LR	SR	IX	CX	SU	X
IR	+	+	+	+	+	+	+	-	-
NR	+	+	+	+	+	+	+	-	-
LR	+	+	+	+	+	+	-	-	-
SR	+	+	+	+	+	-	-	-	-
IX	+	+	+	+	-	+	+	-	-
CX	+	+	+	-	-	+	+	-	-
SU	+	+	+	+	+	-	-	-	-
X	+	-	-	-	-	-	-	-	-

-
-
-

Node Locks

- Node read lock (NR)
 - requires IR locks on the ancestor path
- Level read lock (LR)
 - requested for reading the context node and all nodes located at the level below (all direct-child nodes)
- Child exclusive lock (CX)
 - indicates an X lock on a child
 - defined, in addition to IX, to detect conflicts with LR




Node Locks

Transaction T_1 is reading <price>


Transaction T_2 is reading <book>
(<title>, <author>, and <price>)

Transaction T_3 is modifying
the book title



Node Locks

- Locking subtrees exclusively: intention exclusive lock (IX), child exclusive lock (CX), and exclusive lock (X)
- Exclusive lock (X)
 - requested for updating the context node's content or deleting the context node and its entire subtree
 - requires a CX lock on the parent and IX locks on the ancestors



XML Transactions

Node Locks

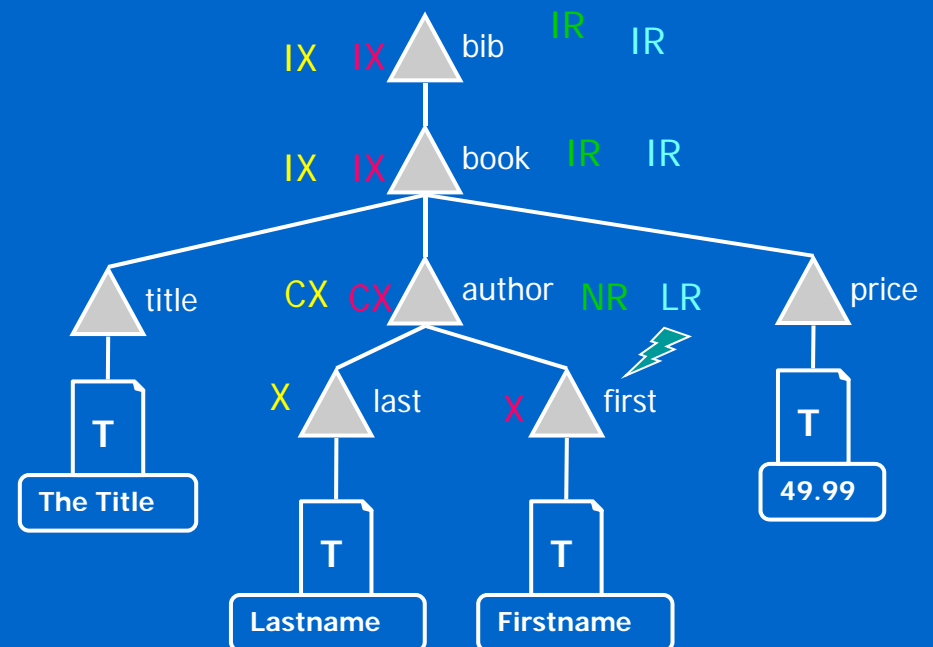
Transaction T_1 is deleting the
<first> node and its content

Transaction T_2 is deleting the
<last> node and its content

Transaction T_3 is reading the
<author> node

Transaction T_4 is reading all
direct-child nodes of <author>


but is blocked when reading
all child nodes of <author> by T_1



-
-
-

Tunable Lock Depth

- Goal
 - reduce the number of locks held by using coarser lock granularity
 - may decrease concurrency
 - when nodes deeper than lock depth are accessed: lock types SR and X are used at the lock depth level



-
-
-

Tunable Lock Depth

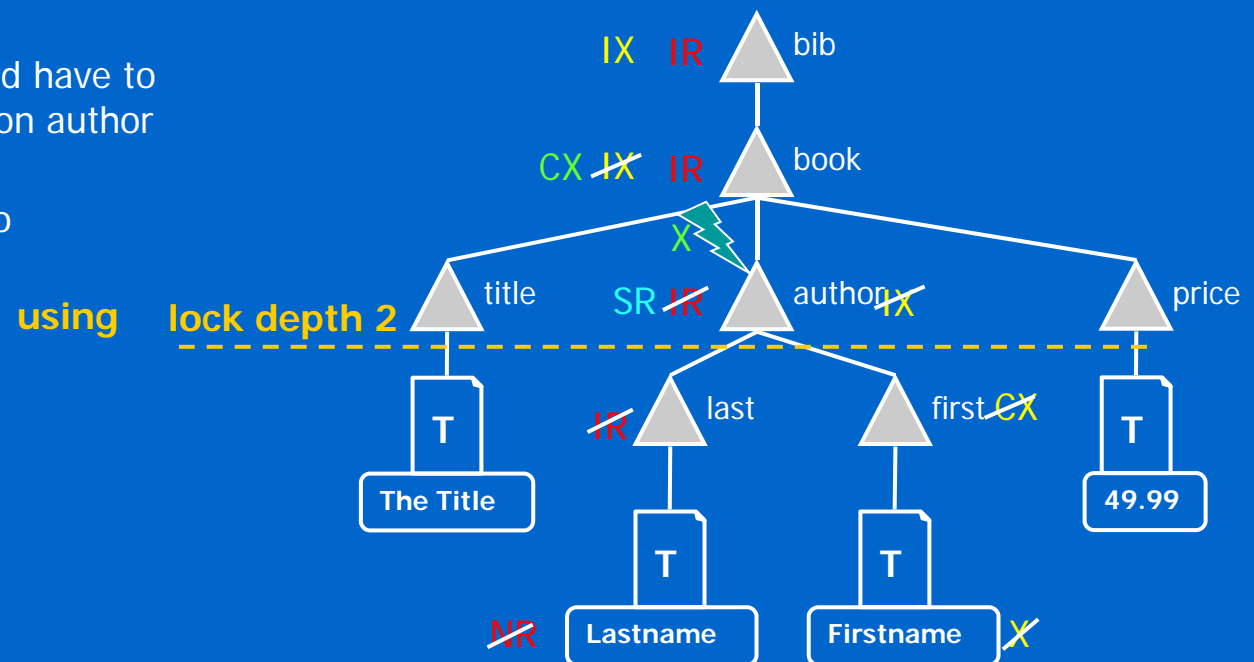
Transaction T_1 is reading
the author's last name

Transaction T_2 is updating
the author's first name

Transaction T_1 would have to
acquire an SR lock on author

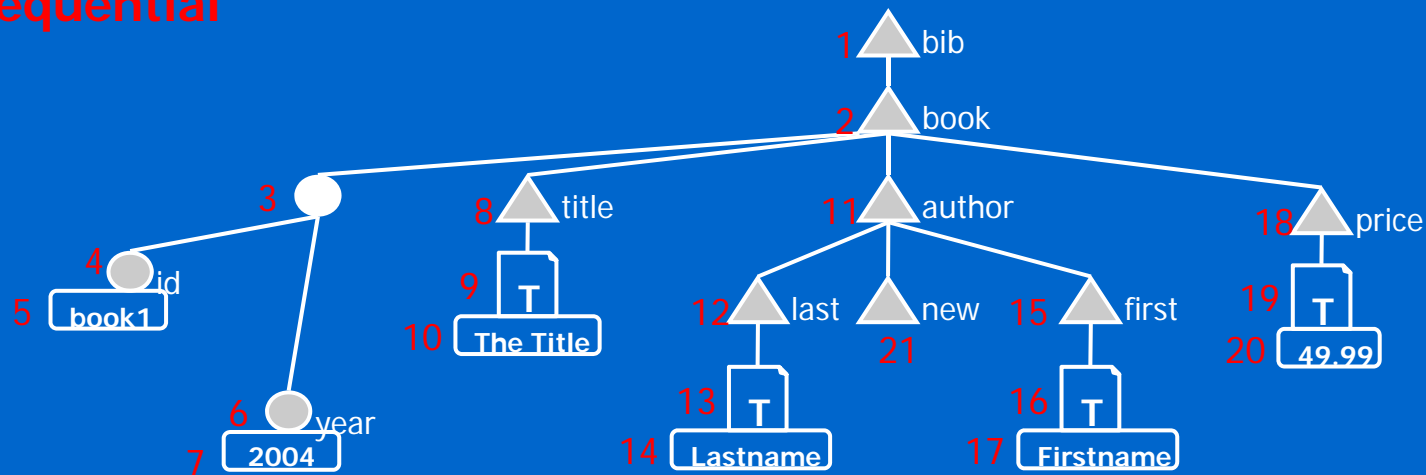
Transaction T_2 would have to
acquire an X lock on author

would therefore have to
wait on author



Identifying Nodes – Node Numbering Schemes

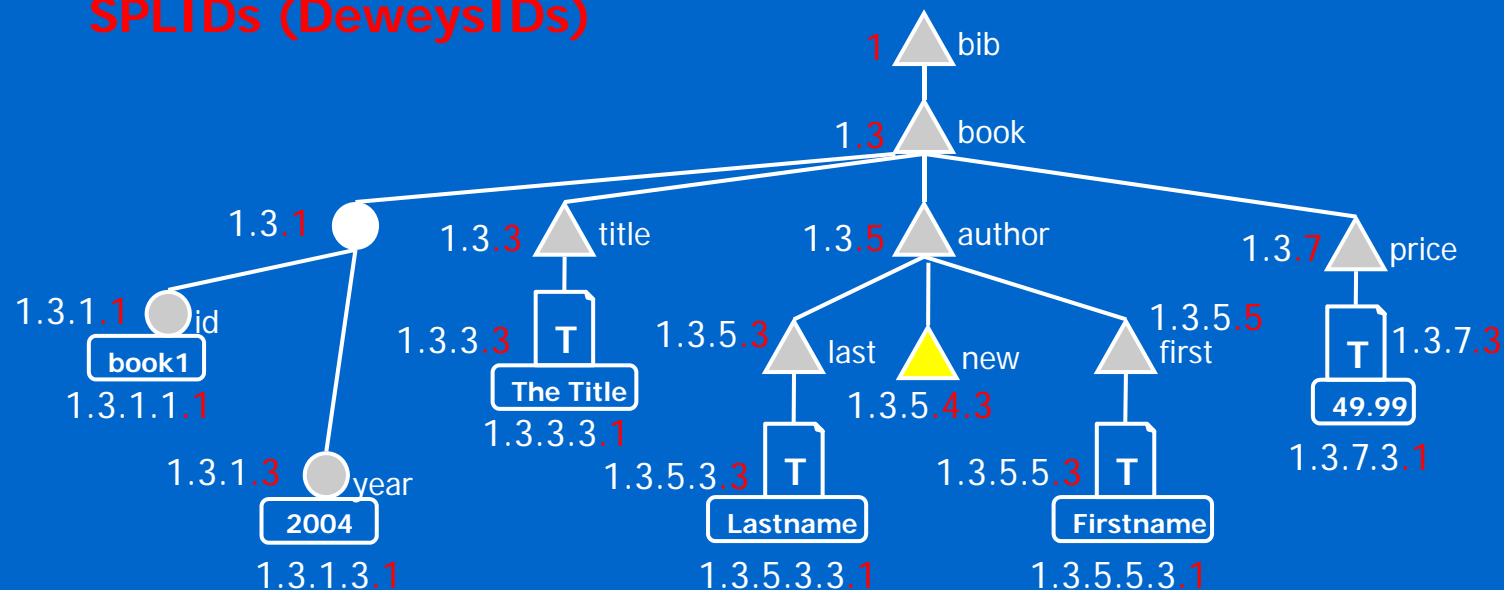
Sequential



- very slow, although supported by on-demand indexes
- determination of parent ID and ancestor IDs, however, is very frequent

Identifying Nodes – Node Numbering Schemes

SPLIDs (DeweysIDs)



-
-
-

OptiX/SnaX

- Locking overhead, especially for read operations, can be tremendous.
- Two snapshot based concurrency control mechanisms that avoid locking.
 - OptiX is a variation of optimistic concurrency control adjusted to use snapshots and work on XML data.
 - SnaX provides the isolation level of snapshot isolation and has similar semantics as the concurrency control mechanisms implemented
- Zeeshan Sardar, Bettina Kemme. “Don’t be a Pessimist: Use Snapshot based Concurrency Control for XML”, ICDE 2006.

-
-
-

Snapshot based Concurrency Control

- Each transaction
 - Working phase, validation phase, update phase
- Two transactions T_i , T_j
 - Concurrent if T_i started the working phase before T_j finished the update phase and committed
- Validation
 - OptiX: WriteSet(T_j) not overlaps with ReadSet(T_i)
 - SnaX: WriteSet(T_j) not overlaps with WriteSet(T_i)
 - Snapshot isolation as in many RDBMS
 - Assuming that many applications are read only

-
-
-

ReadSets

- XML structure
 - $RR(T_i)$: roots of subtree returned as part of the query result
 - $ER(T_i)$: nodes that are explicitly read as part of a predicate or path constraint, but not return as part of result
 - $ERa(T_i)$: nodes that are read for the insertion of a sibling after them

-
-
-

WriteSets

- XML structure
 - $D(T_i)$: roots of subtree deleted or replaced
 - $R_n(T_i)$: nodes to be renamed
 - $I(T_i)$: the immediate parents of any nodes inserted by T_i
 - $I_a(T_i)$: the same set of nodes as in $Era(T_i)$

-
-
-

OptiX Validation

- Suppose T_i wants to validate, and concurrent transaction T_j validated before T_i .
- When $\text{ReadSet}(T_i)$ overlaps with $\text{WriteSet}(T_j)$
 - Instead of immediately inducing a conflict, we look at the subsets of $\text{ReadSet}(T_i)$ and $\text{WriteSet}(T_j)$ of which p is a member.
 - YES indicates compatibility, while NO shows a conflict leading to an abort of T_i .
- A conflict means that if T_i had executed serially after T_j , then it would have possibly read a different value for p than in the concurrent execution.
- In the table, q is a descendant of p .

-
-
-

OptiX Validation

Tj already validated on p

Ti validating	D	Rn	I	la
On p RR	No	No	No	Yes
On p ER	No	No	Yes	Yes
On p ERa	No	No	Yes	No
On q ReadSet	No	Yes	Yes	Yes

-
-
-

SnaX Validation

- In SnaX, remember snapshot isolation does not guarantee that there is an equivalent serial execution.
 - ignore the read sets and only consider write/write conflicts.
- Generally, delete and replace conflict with most other update types, while inserts conflict only with few.

-
-
-

SnaX Validation

Tj already validated on p

Ti validating	D	Rn	I	la
On p D	No	No	No	Yes
On p Rn	No	No	Yes	Yes
On p I	No	Yes	Yes	Yes
On p la	Yes	Yes	Yes	No
On q WriteSet	No	Yes	Yes	Yes