

Generating **Efficient** Execution Plans for Vertically Partitioned XML Databases

Research paper review by

QING Pei, Edward	11500811g
LO Wing Yi, Wing	11523479g
SHAO Shuai, Philip	11552402g

April 10, 2012

What ?

Why ?


How ?

What ?

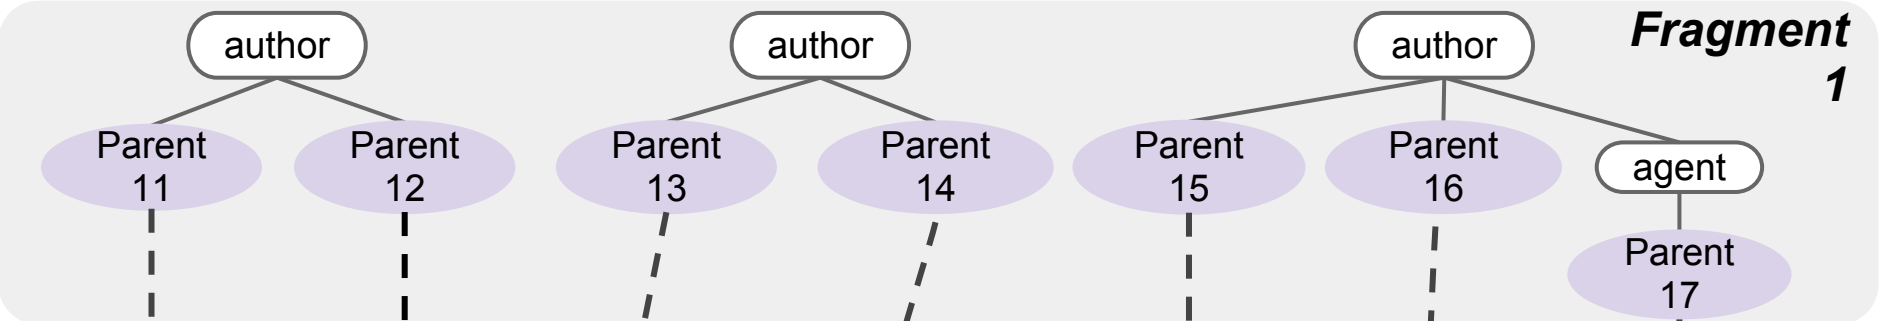
Query Processing

	Centralized	Distributed
RDBMS	✓	✓
XML	✓	<i>This paper</i>

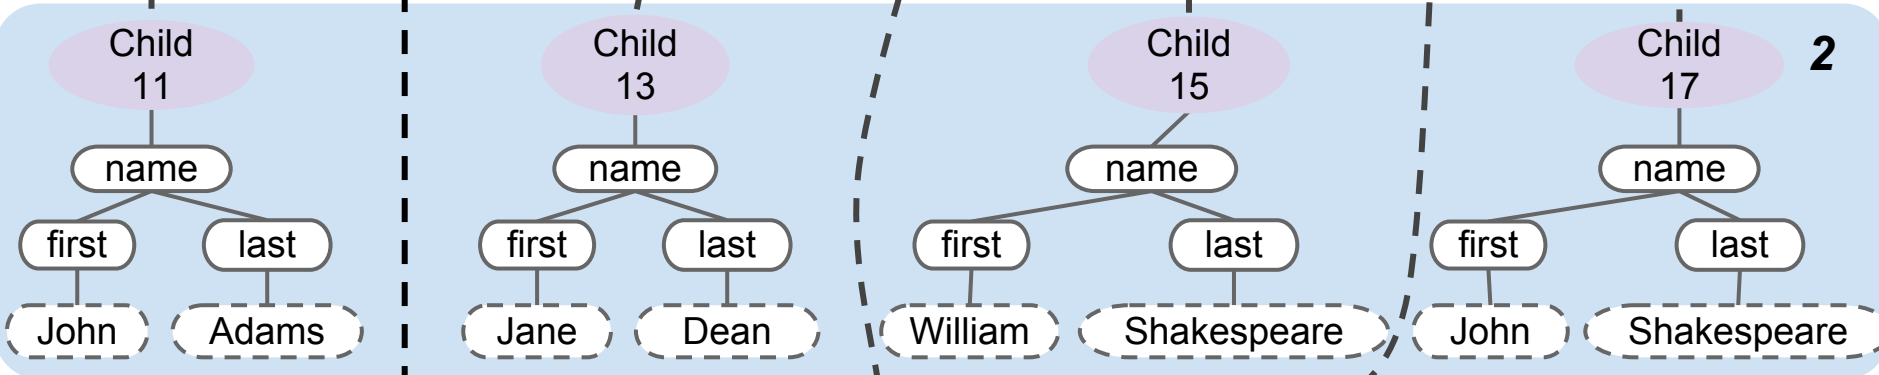
XML in the Cloud



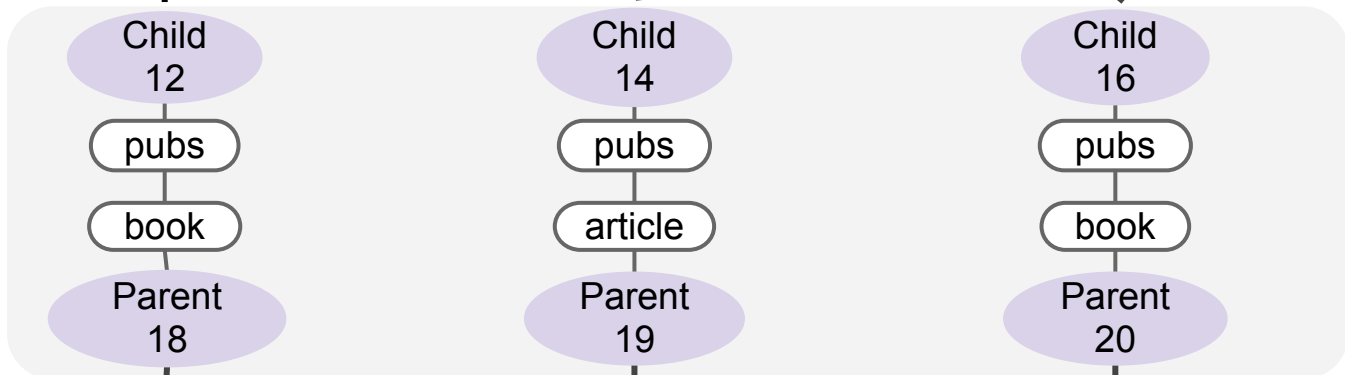
Fragment 1



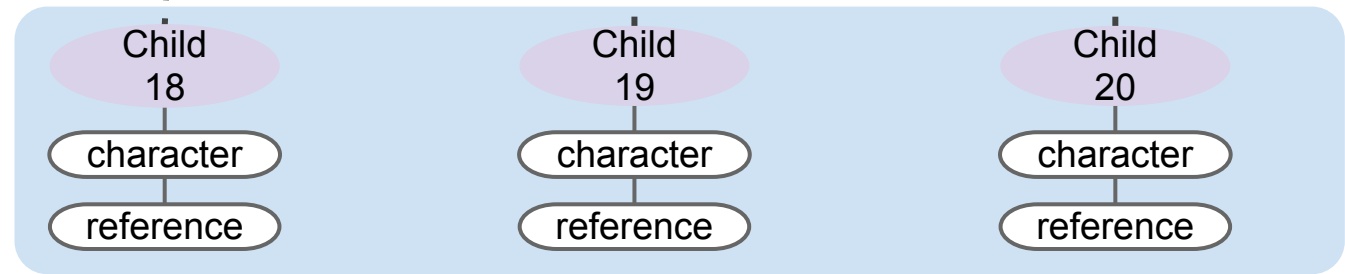
2



3



4



Why ?

Distributed architecture leads to **Different** execution plans

For a single query, the **order** in which *joins* are performed results in various time consumed.

$$\begin{aligned} &\text{Response time} \\ &= \\ &\text{local execution time} \\ &+ \\ &\text{joining time} \end{aligned}$$

local execution time

$snip(i)$: the number of document subtrees
accessed by the local plan at *fragment i*

smaller $snip(i)$ preferred

joining time

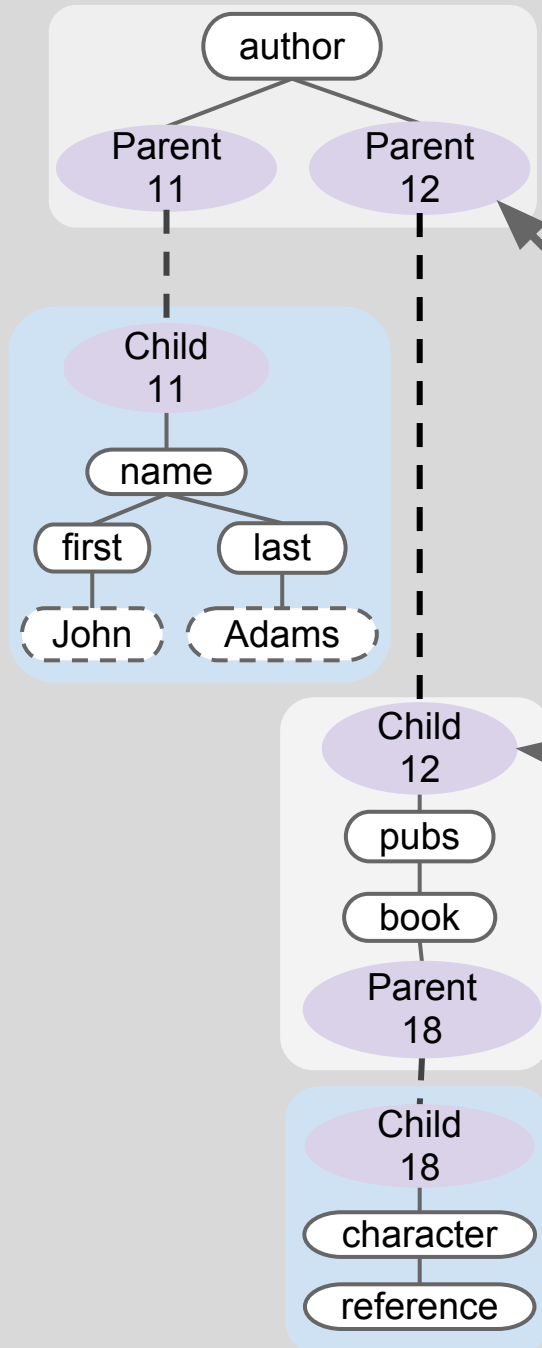
card(i): the number of tuples that are returned by the local plan when evaluated at *fragment i*

smaller *card(i)* preferred

***Which* plan has the
minimum response time?**

How ?

**keep the relation
between
fragments**



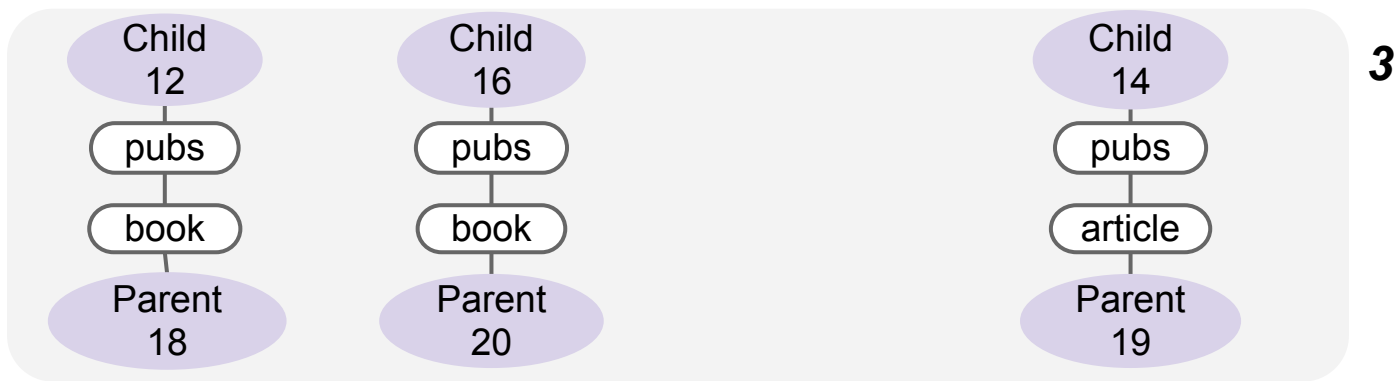
proxy nodes

Optimizing distributed plans




Optimizing distributed plans

Pushing Cross-Fragment Joins


fully works on left-deep plans



re-partition



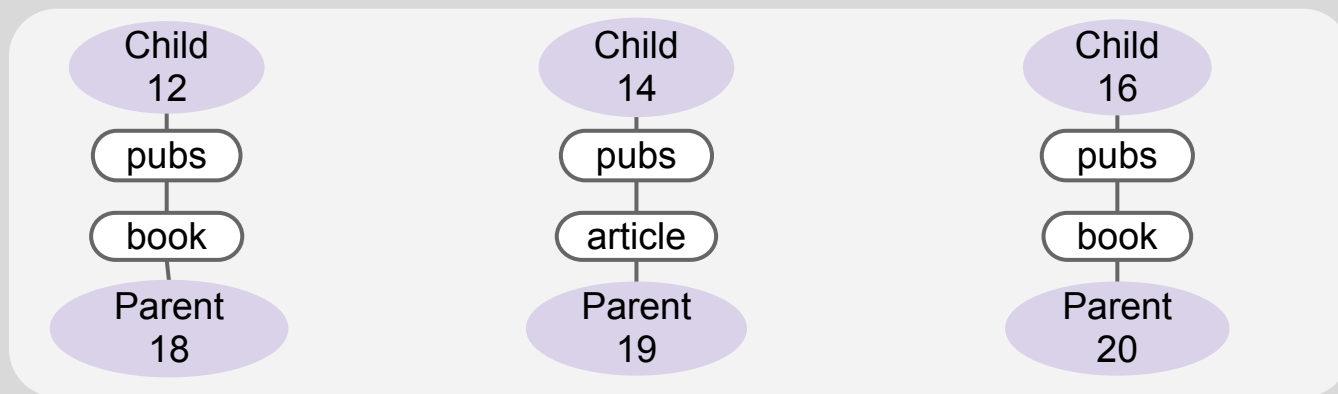
only visit children in 4 with parents in 3.2)



Optimizing distributed plans


Label Path Filtering

//**book**//reference




Evaluation

Centralized & Distributed Techniques Comparison (Collection 3.5GB)



Selective XPathMark Performance Results (Collection 12GB)



Conclusion

Greatly improves response time
of querying large XML collections.



Small overhead. Choosing the
fastest plan took **< 0.01** seconds.

Q & A

**Merci
beaucoup**

Appendix


Distributed Execution Plans



left-deep execution plans

Appendix

Distributed Execution Plans



not a left-deep execution plan

Appendix

Queries used for evaluation

Q1 /open auction[initial > 200]//item//mail/from

Q2 /open auction[initial > 200][.//author/person/name[starts-with(., 'Ry')]]//item//mail/from

Q3 /open auction[initial > 200][.//author/person/name[starts-with(., 'Ry')]]//item//category/id

Q4 /open auction[initial > 200][.//author/person[profile/age > 30]/name[starts-with(., 'Ry')]]//item//category/id

Q5 /open auction[initial > 200]//author/person[starts-with(name, 'Ry')]/profile/interest/category/description

Appendix

Queries used for XPathMark

- A1** /site/closed auctions/closed auction/annotation/description/text/keyword
- A2** //closed auction//keyword
- A3** /site/closed auctions/closed auction//keyword
- A4** /site/closed auctions/closed auction [annotation/description/text/keyword]/date
- A5** /site/closed auctions/closed auction[descendant:: keyword]/date
- A6** /site/people/person[profile/gender and profile/age]/name
- B7** //person[profile/@income]/name

Appendix

Queries used for Selective XPathMark

A1S /site/closed auctions/closed auction[price > 600]
/annotation/description/text/keyword

A2S //closed auction[price > 600]//keyword

A3S /site/closed auctions/closed auction[price > 600]
//keyword

A4S /site/closed auctions/closed auction[price > 600]
[annotation/description/text/keyword]/date

A5S /site/closed auctions/closed auction[price > 600]
[descendant::keyword]/date

A6S /site/people/person[starts-with(name, 'Ry')]
[profile/gender and profile/age]/name

B7S //person[starts-with(name, 'Ry')][profile/@income]/name

Appendix

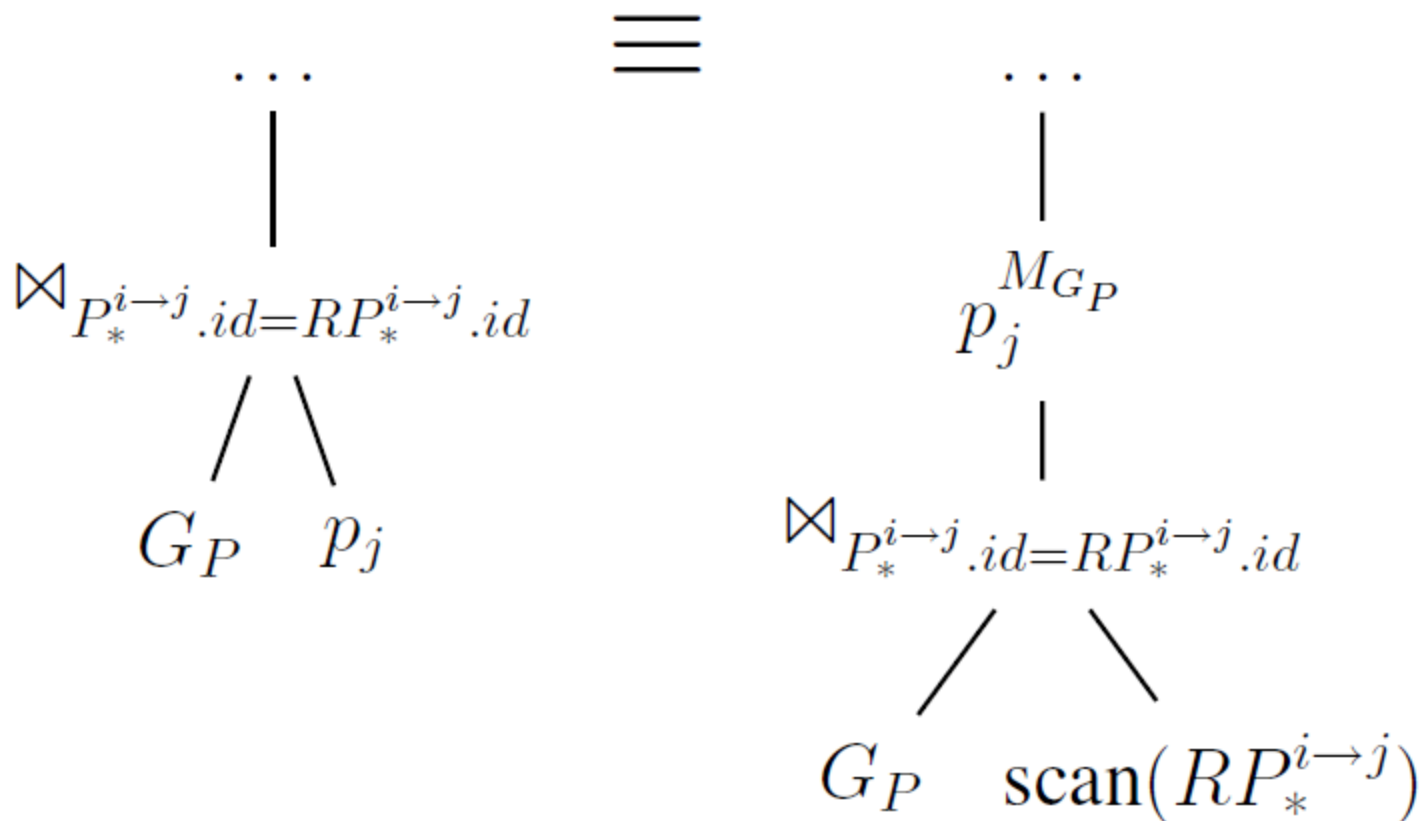


Figure 11: Cross-fragment join pushing rewrite

Appendix

$$\begin{array}{c} p_j \quad \equiv \quad p'_j \\ | \\ \sigma_{RP_*^{i \rightarrow j}.label \in L_j} \\ | \\ \text{scan}(RP_*^{i \rightarrow j}) \end{array}$$

Figure 12: Label path rewrite

Appendix

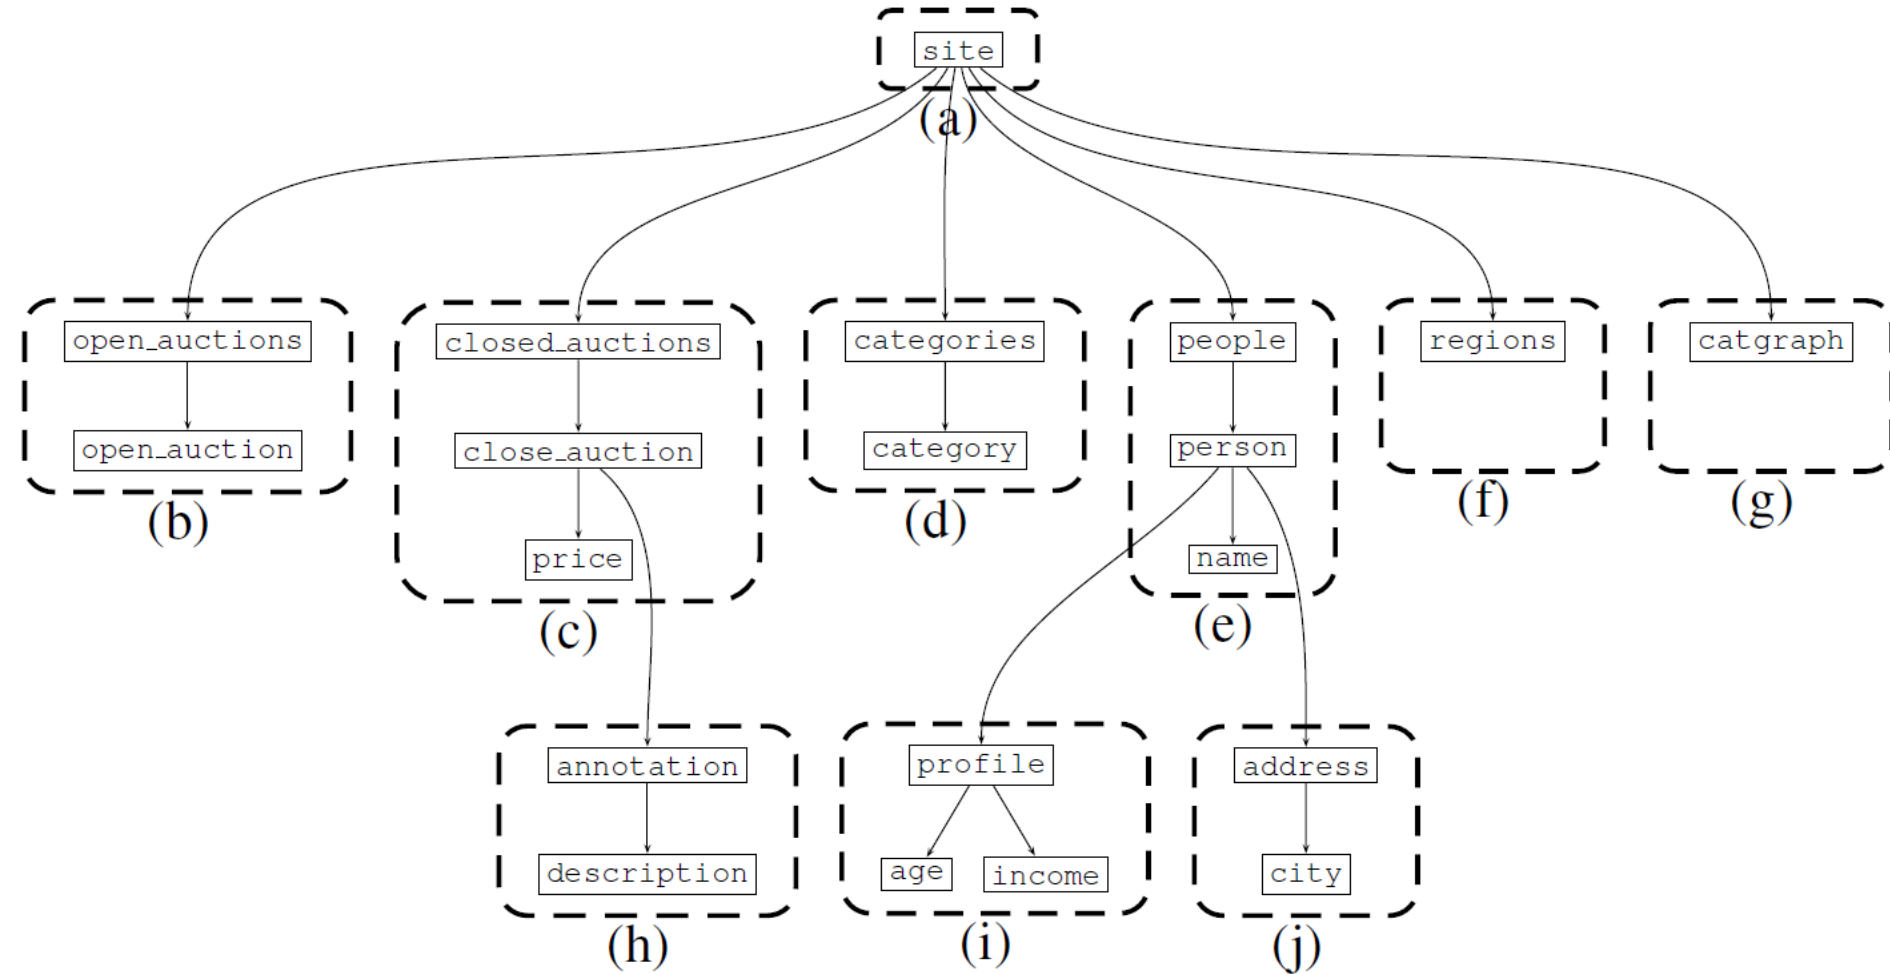


Figure 13: Fragmentation schema used in second experiment

Appendix

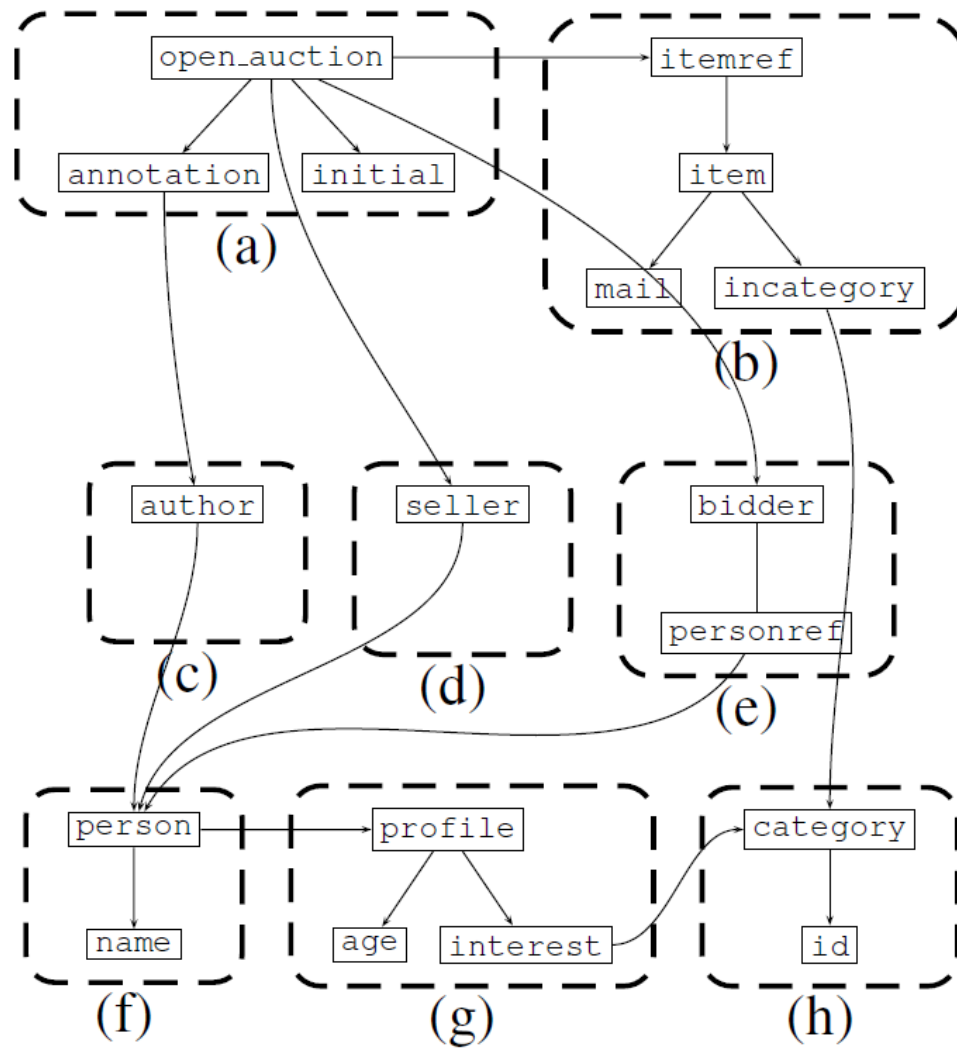


Figure 14: Fragmentation schema used in first experiment