

Generating **Efficient** Execution Plans for Vertically Partitioned XML Databases

Research paper review by

QING Pei, Edward 11500811g
LO Wing Yi, Wing 11523479g
SHAO Shuai, Philip 11552402g

April 10, 2012

What ?

Why ?

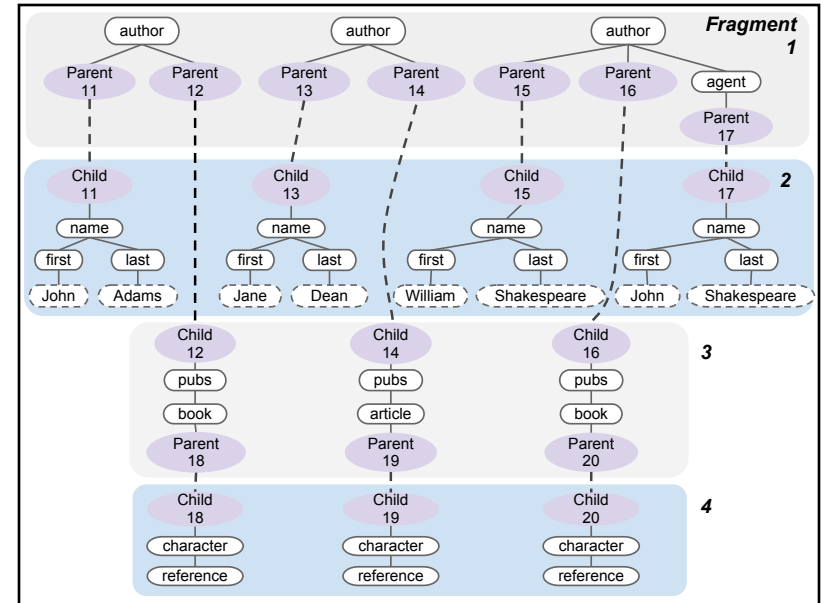

How ?

What ?

Query Processing

	Centralized	Distributed
RDBMS	✓	✓
XML	✓	<i>This paper</i>

XML in the Cloud



Why ?

Distributed architecture
leads to
Different execution plans

For a single query, the **order** in which *joins* are performed results in various time consumed.

$$\begin{aligned} \text{Response time} \\ = \\ \text{local execution time} \\ + \\ \text{joining time} \end{aligned}$$

local execution time

snip(i): the number of document subtrees accessed by the local plan at *fragment i*

smaller *snip(i)* preferred


joining time

card(i): the number of tuples that are returned by the local plan when evaluated at *fragment i*

smaller *card(i)* preferred

***Which* plan has the
minimum response time?**

How ?

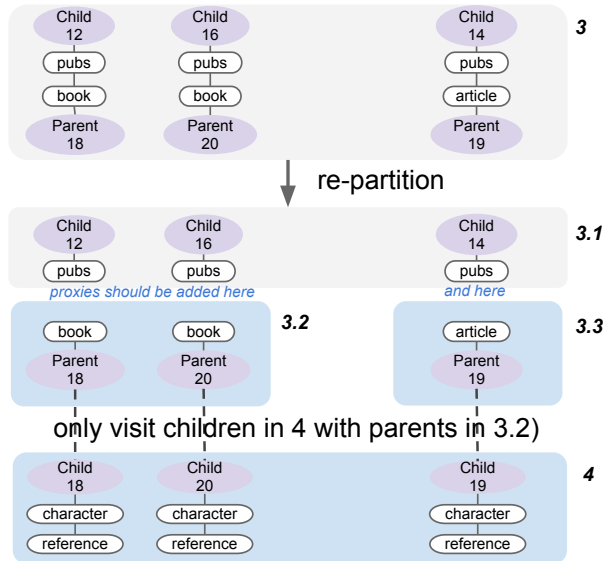


Optimizing distributed plans

Optimizing distributed plans

Pushing Cross-Fragment Joins


fully works on left-deep plans



Optimizing distributed plans


Label Path Filtering

//book//reference

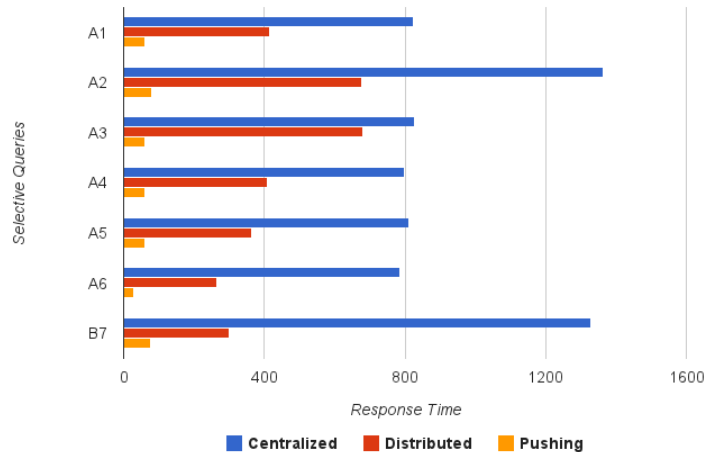


Evaluation

Centralized & Distributed Techniques Comparison (Collection 3.5GB)



**Selective XPathMark Performance Results
(Collection 12GB)**



Conclusion

Greatly improves response time of querying large XML collections.

Small overhead. Choosing the fastest plan took **< 0.01** seconds.

Q & A