

# Generating **Efficient** Execution Plans for Vertically Partitioned **XML** Databases

Research paper review by

QING Pei, Edward	11500811g
LO Wing Yi, Wing	11523479g
SHAO Shuai, Philip	11552402g

*April 10, 2012*

**What ?**

**Why ?**

**How ?**

For this paper, as well as today's presentation, we raised three questions.

# What ?

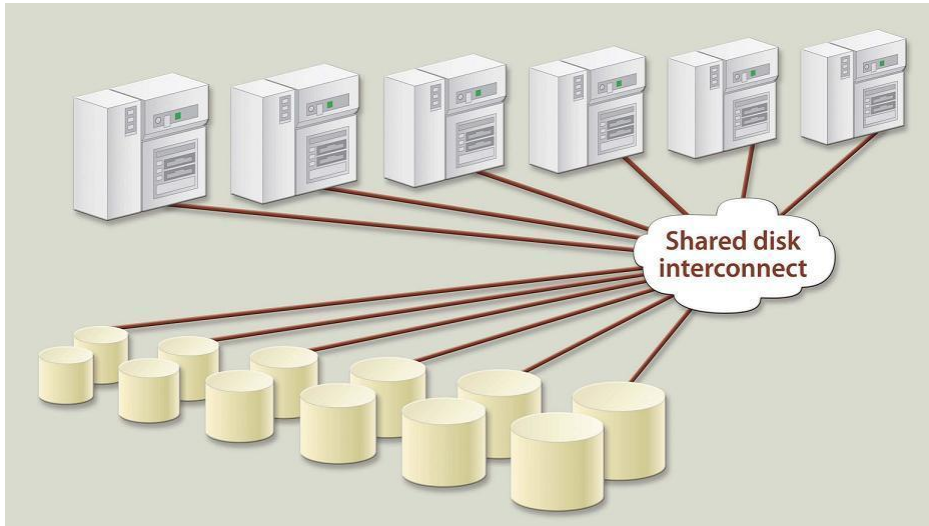
What is this paper talking about?

## Query Processing

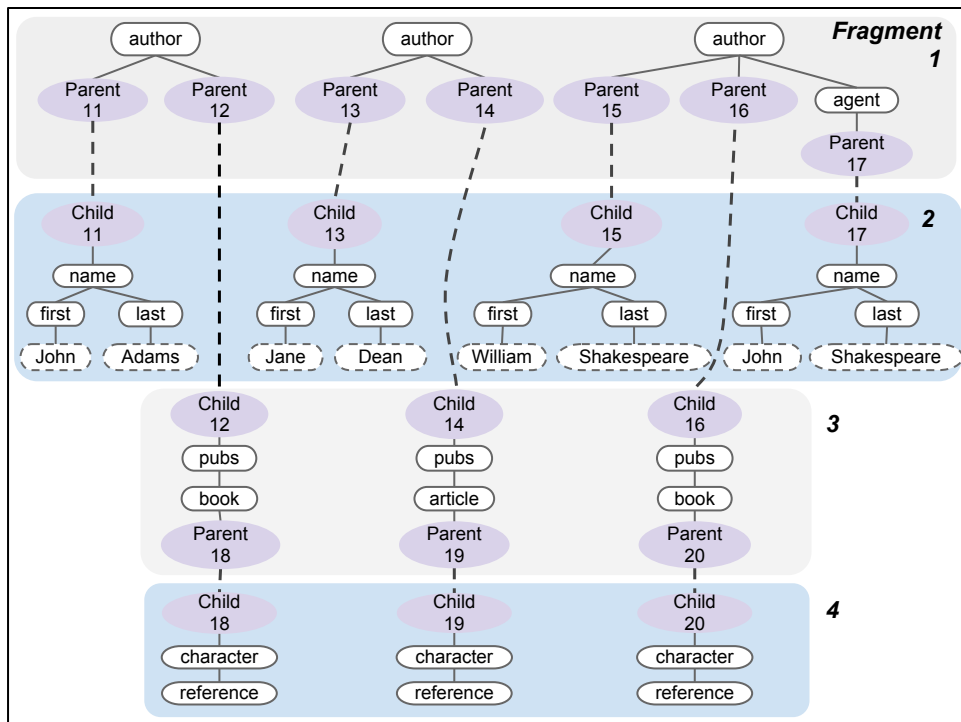
	Centralized	<b>Distributed</b>
RDBMS	✓	✓
<b>XML</b>	✓	<i><b>This paper</b></i>

It is focused on the query processing of distributed XML databases, other than a centralized one or relational databases.

# XML in the Cloud



Several XML servers instead of a centralized one are distributed in the cloud



As shown in this picture, all data are originally in one tree.

For better scalability, it has been vertically partitioned into four fragment.

So there are four (maybe five) distributed server here, restoring local data respectively.

That is what this paper is going to talk about.

# Why ?

Now we come to the second question, why. Why this paper exist?

**Distributed** architecture  
leads to  
**Different** execution plans

For a single query, the **order** in which *joins* are performed results in various time consumed.

For a single query, there may exist different plans for local execution/joining operations.



$$\begin{aligned} &\textbf{Response time} \\ &= \\ &\textbf{local execution time} \\ &+ \\ &\textbf{joining time} \end{aligned}$$

The equation

In naive distributed xml database query execution, if we want to execute a query, we have to access every tuple and every node on each distributed server. This is quite slow.

The join operation is also slow. Then it leads to even longer response time.

That's why we have this paper - to reduce the response time.

## local execution time

*snip(i)*: the number of document subtrees  
accessed by the local plan at *fragment i*

smaller *snip(i)* preferred

local execution time depends on snip

## joining time

*card(i)*: the number of tuples that are returned by the local plan when evaluated at *fragment i*

smaller *card(i)* preferred

joining time depends on card

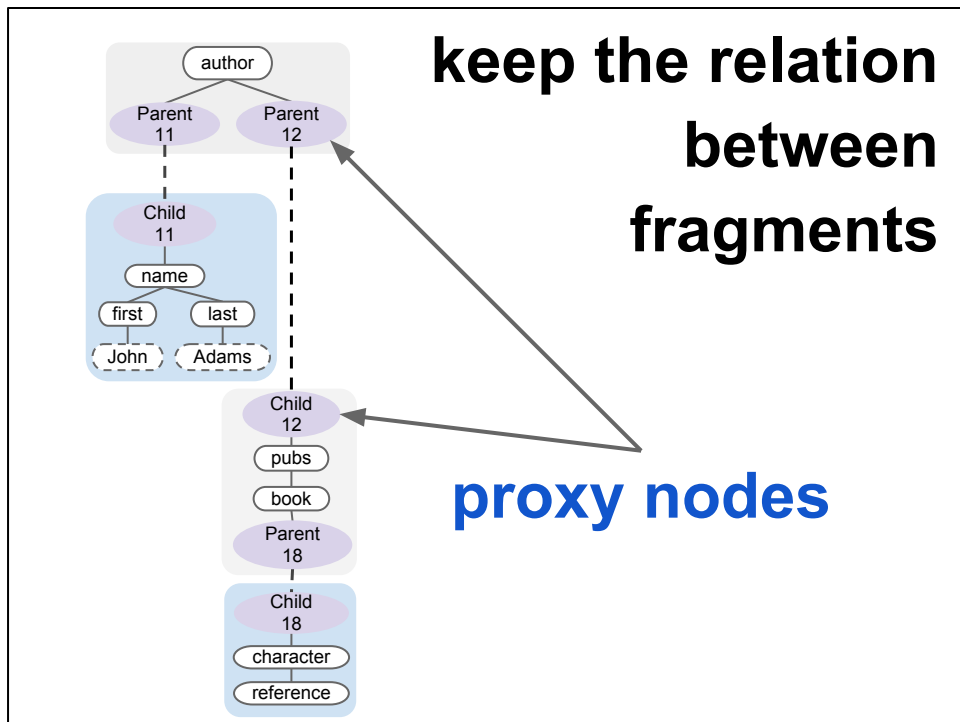
***Which* plan has the  
*minimum* response time?**

For the distributed plan, there are several optimizations for it.

This paper is trying to find fastest one.

# How ?

Now, the third question, how.  
How is it done?



One subtree of the original tree. Author and name elements are connected.

Some edges are broken during partitioning.

Proxy nodes are then introduced to recover the relations.

# **Optimizing distributed plans**

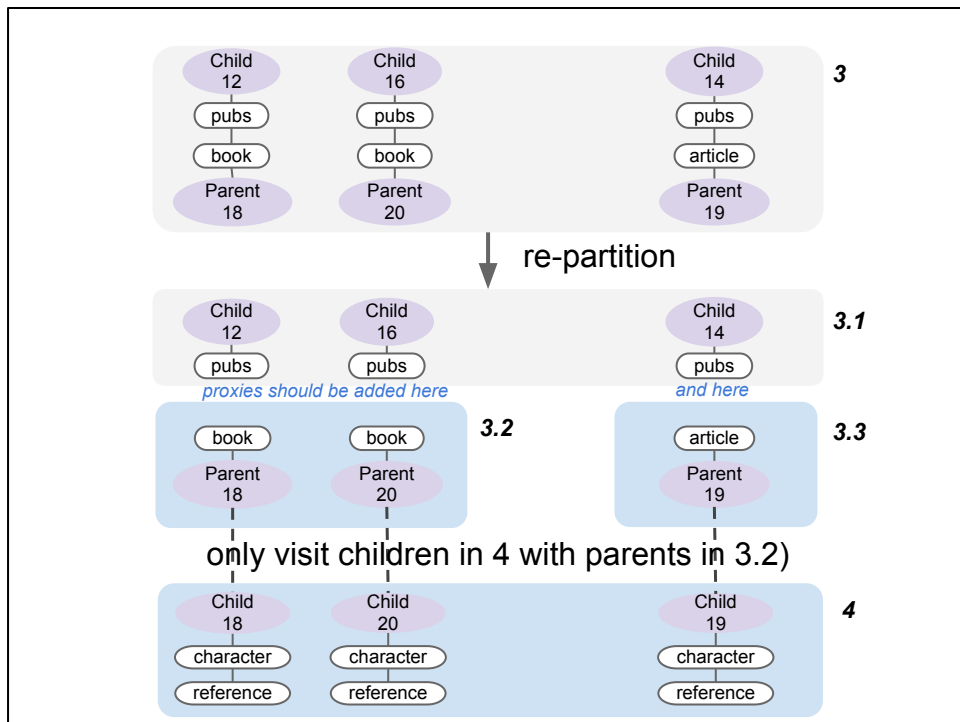
There are several optimization plans, and this paper talks about two of them.

Optimizing distributed plans

# Pushing Cross-Fragment Joins

fully works on left-deep plans





To illustrate how this works, we modify the example a little bit.

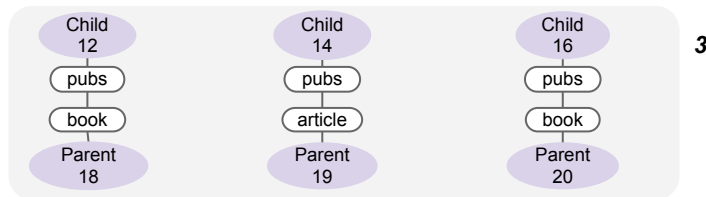
In the original partition, fragment three has been re-partitioned into two fragments so that the proxy nodes could work.

The proxy nodes reduce visits to unnecessary subtrees.

## Optimizing distributed plans

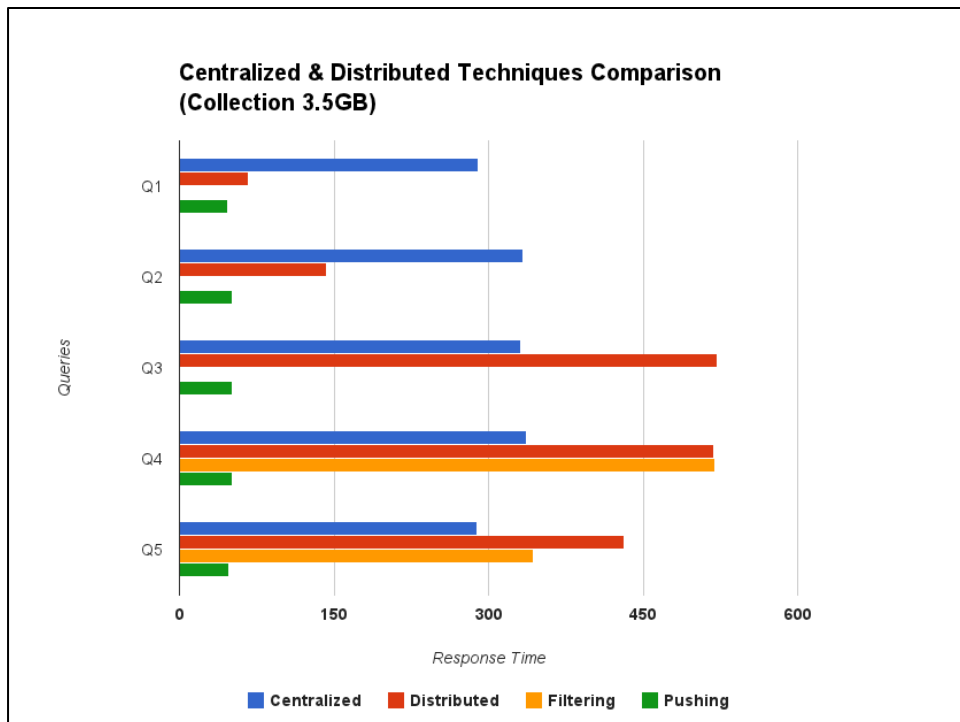
# Label Path Filtering

//**book**//reference

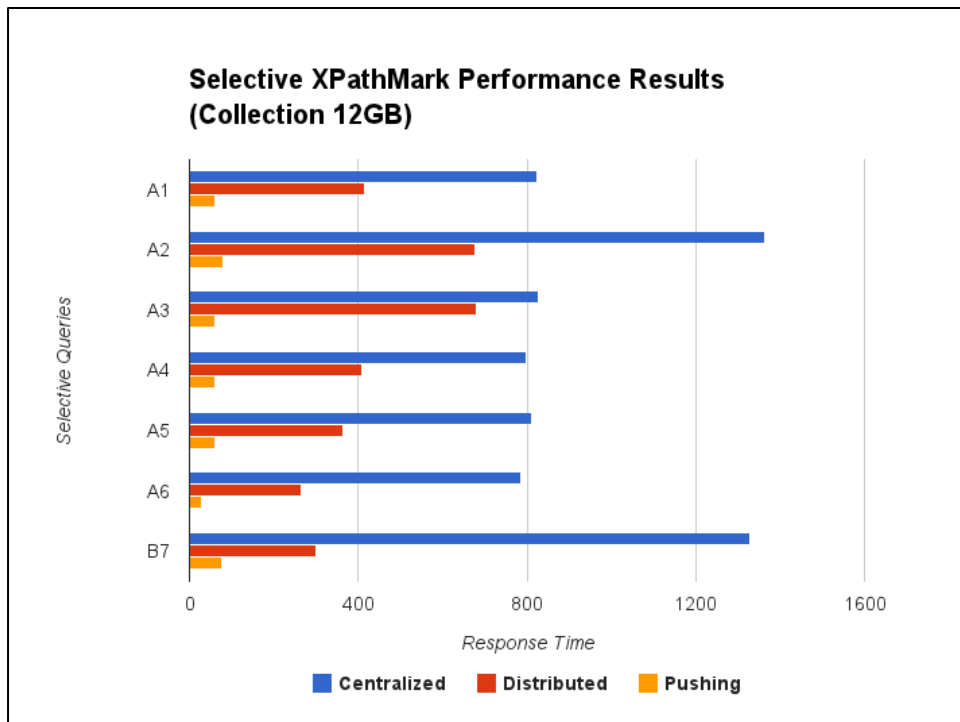


works in the situation that when a parent has several children that share its grandchild  
It reduced the response time, snip, by avoiding access to some unnecessary subtree

# Evaluation



We could see that Label Path filtering has its restrictions -  
it works in the situation that when a parent has several children that share its grandchild



For selective XPath queries, the proposed optimizations greatly reduces the response time. It can get up to 10 times faster execution.

# Conclusion

**Greatly** improves response time of querying large XML collections.

Small overhead. Choosing the fastest plan took **< 0.01** seconds.

Scalability.

**Q & A**

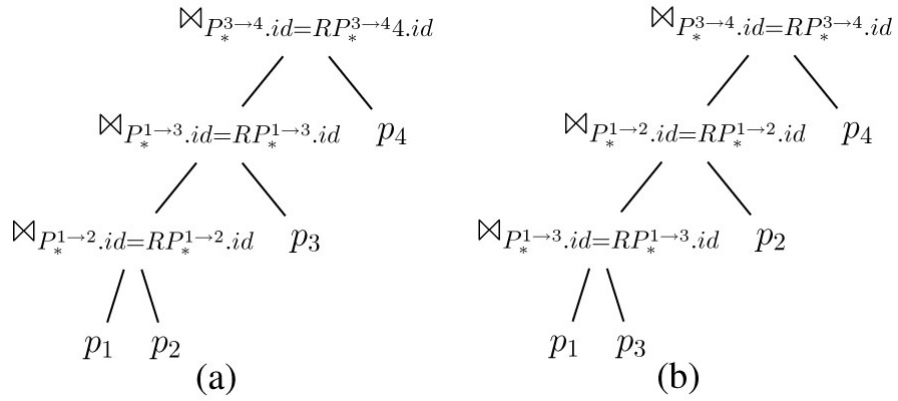


**Merci  
beaucoup**

The end.

## Appendix

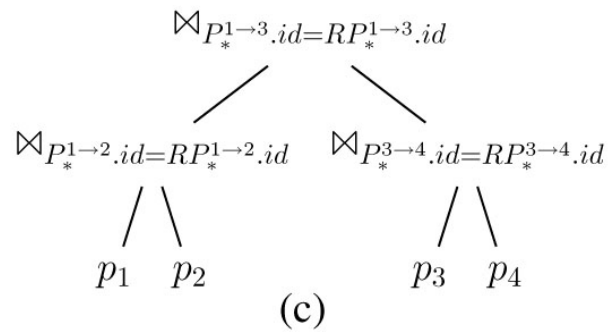
### Distributed Execution Plans



left-deep execution plans

## Appendix

### Distributed Execution Plans



not a left-deep execution plan

## Appendix

### Queries used for evaluation

**Q1** /open auction[initial > 200]//item//mail/from

**Q2** /open auction[initial > 200][.//author/person/name[starts-with(., 'Ry')]]//item//mail/from

**Q3** /open auction[initial > 200][.//author/person/name[starts-with(., 'Ry')]]//item//category/id

**Q4** /open auction[initial > 200][.//author/person[profile/age > 30]/name[starts-with(., 'Ry')]]//item//category/id

**Q5** /open auction[initial > 200]//author/person[starts-with(name, 'Ry')]/profile/interest/category/description

## Appendix

### Queries used for XPathMark

**A1** /site/closed auctions/closed  
auction/annotation/description/text/keyword

**A2** //closed auction//keyword

**A3** /site/closed auctions/closed auction//keyword

**A4** /site/closed auctions/closed auction  
[annotation/description/text/keyword]/date

**A5** /site/closed auctions/closed auction[descendant::  
keyword]/date

**A6** /site/people/person[profile/gender and profile/age]/name

**B7** //person[profile/@income]/name

## Appendix

### Queries used for Selective XPathMark

**A1S** /site/closed auctions/closed auction[price > 600]  
/annotation/description/text/keyword

**A2S** //closed auction[price > 600]//keyword

**A3S** /site/closed auctions/closed auction[price > 600]  
//keyword

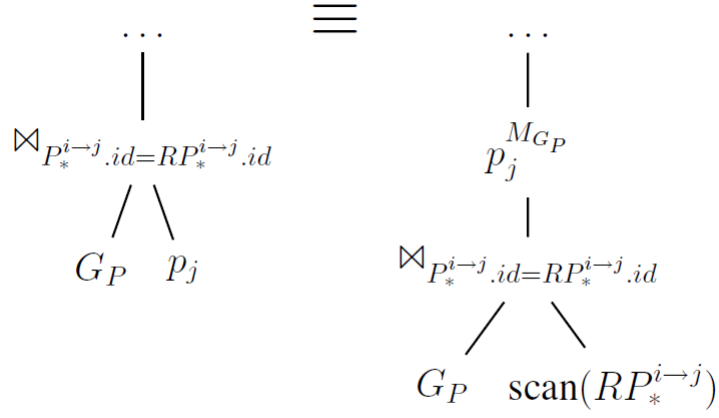
**A4S** /site/closed auctions/closed auction[price > 600]  
[annotation/description/text/keyword]/date

**A5S** /site/closed auctions/closed auction[price > 600]  
[descendant::keyword]/date

**A6S** /site/people/person[starts-with(name, 'Ry')]  
[profile/gender and profile/age]/name

**B7S** //person[starts-with(name, 'Ry')][profile/@income]/name

## Appendix



**Figure 11: Cross-fragment join pushing rewrite**

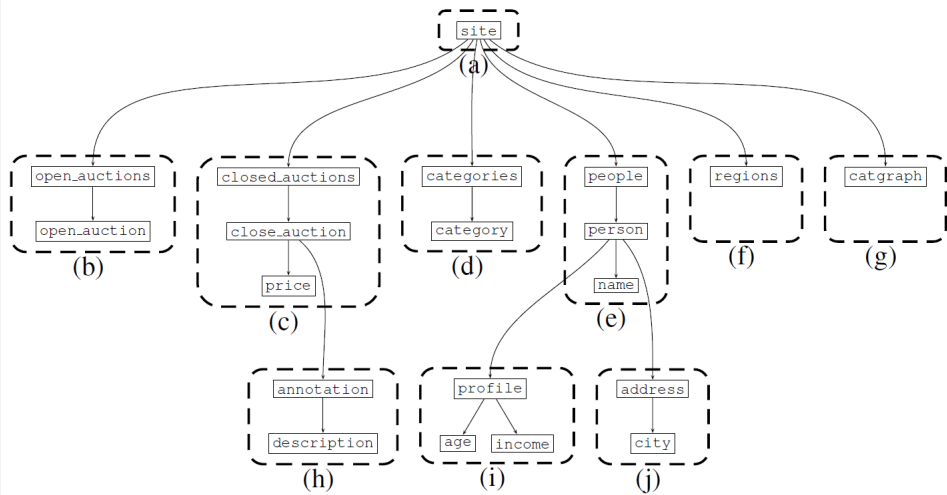
## Appendix

$$\begin{array}{ccc} p_j & \equiv & p'_j \\ & & | \\ & & \sigma_{RP_*^{i \rightarrow j}.label \in L_j} \\ & & | \\ & & \text{scan}(RP_*^{i \rightarrow j}) \end{array}$$

**Figure 12: Label path rewrite**

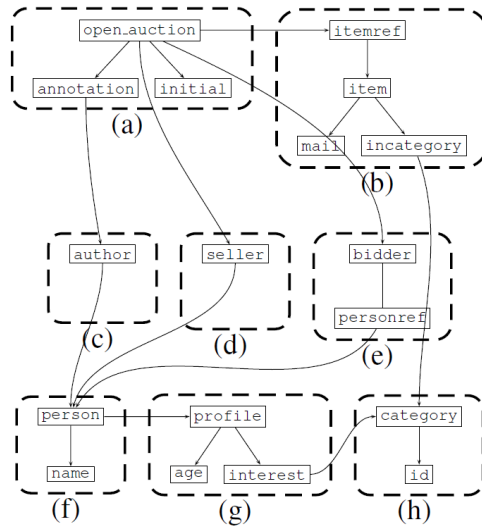


## Appendix



**Figure 13: Fragmentation schema used in second experiment**

## Appendix



**Figure 14: Fragmentation schema used in first experiment**